


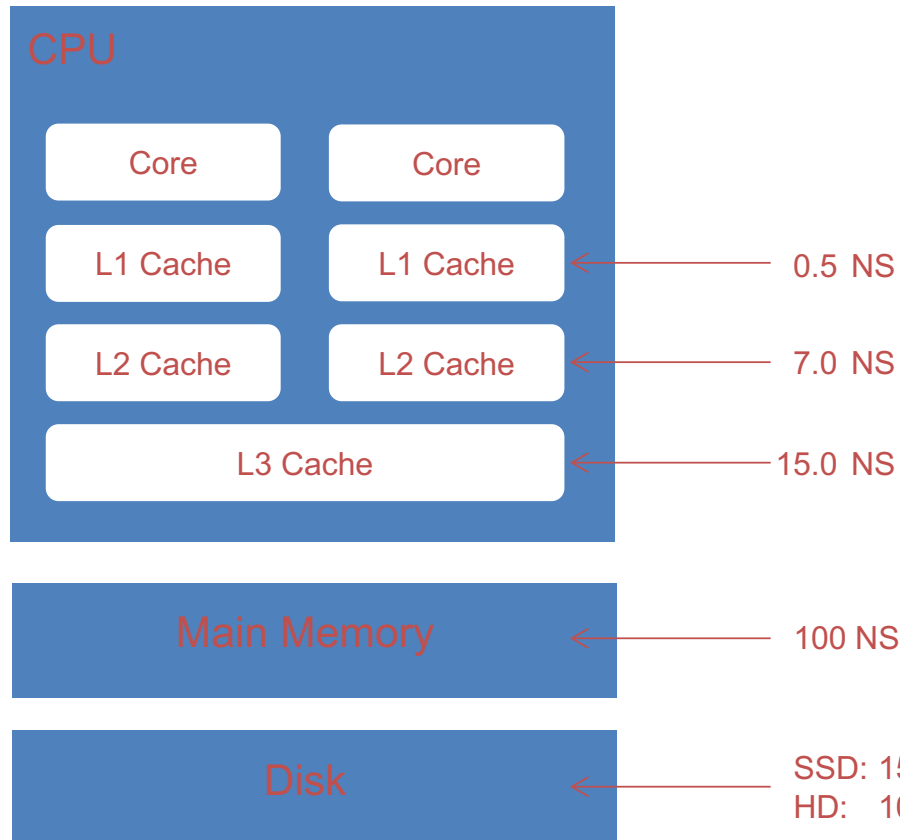
# The Rise of RAMCloud

Rong Chen

# Outline

- Overview
- Motivation
- Durability
- Crash Recovery

# Review: Latency Numbers



Yes, DRAM is 100,000 times faster than disk, but DRAM access is still 6-200 times slower than on-chip caches

# RAMCloud Overview

Storage for datacenters  
1000-10000 commodity  
servers

64 GB DRAM/server

All data always in RAM

Durable and available


Performance goals:

- High throughput:

- 1M ops/sec/server

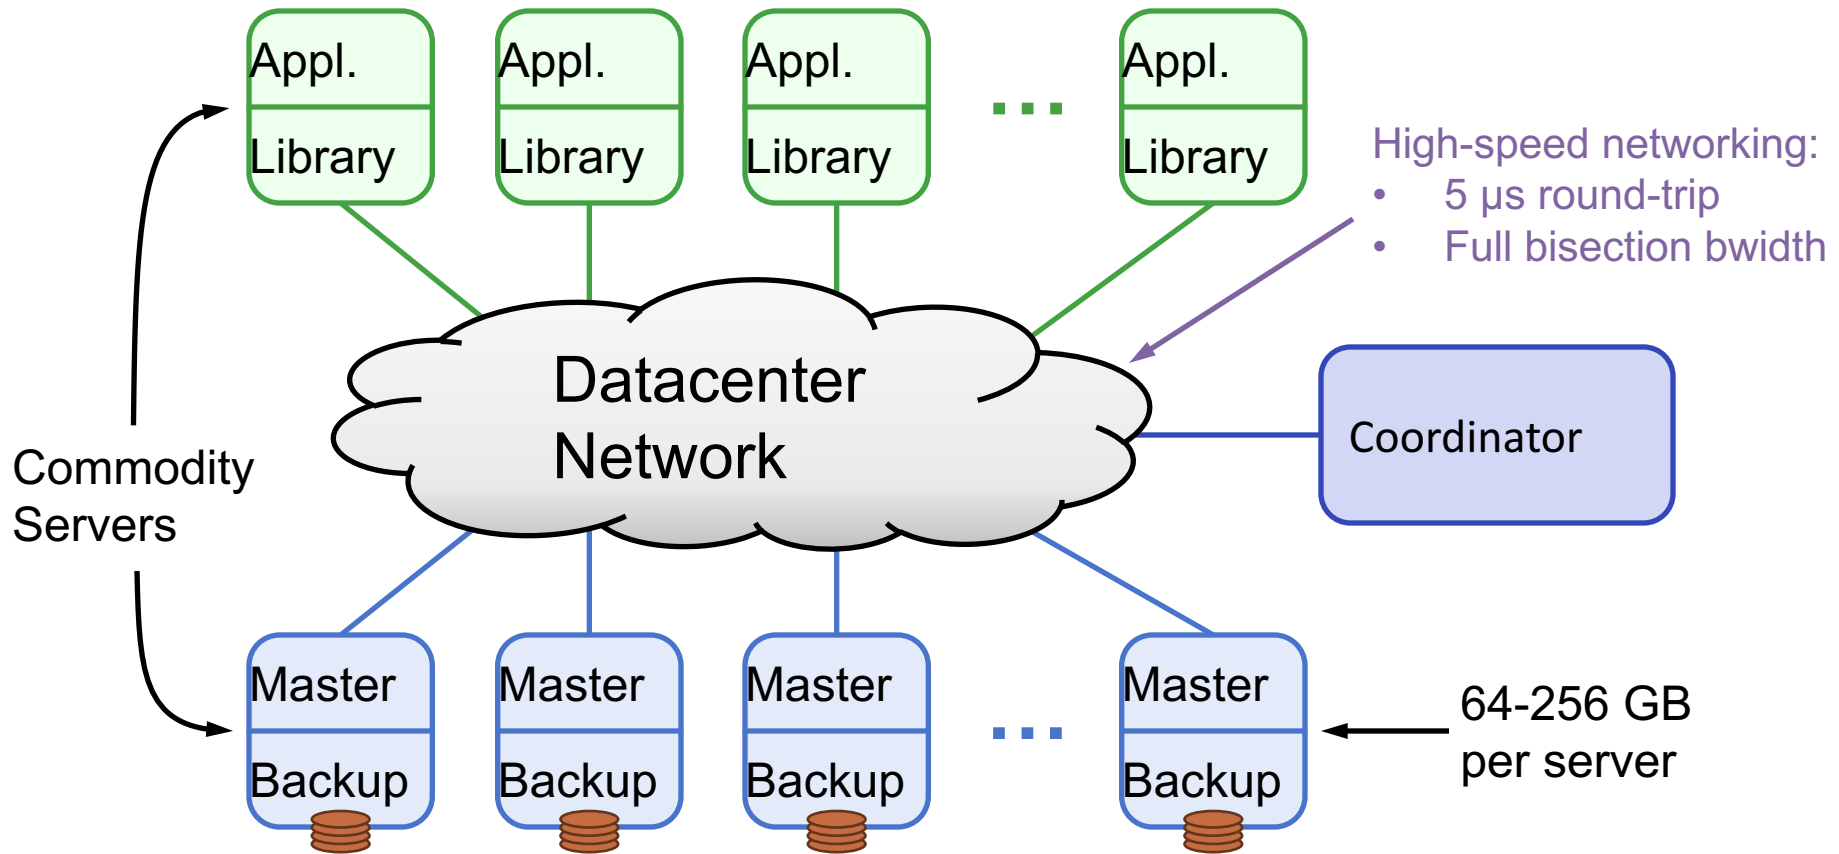
- Low-latency access:

- 5-10 $\mu$ s RPC



# RAMCloud Architecture

**1000 – 100,000 Application Servers**



**1000 – 10,000 Storage Servers**

# Data Model: Key-Value Store

- Basic operations:

- `read(tableId, key)`  
=> blob, version
- `write(tableId, key, blob)`  
=> version
- `delete(tableId, key)`

(Only overwrite if  
version matches)



- Other operations:

- `cwrite(tableId, key, blob, version)`  
=> version
- Enumerate objects in table
- Efficient multi-read, multi-write
- Atomic increment

- Not in RAMCloud 1.0:

- Atomic updates of multiple objects
- Secondary indexes

## Tables



Object

Key ( $\leq 64\text{KB}$ )

Version (64b)


Blob ( $\leq 1\text{MB}$ )

# Example Configurations

	2010	5-10 years
<b># servers</b>	<b>1000</b>	<b>1000</b>
<b>GB/server</b>	<b>64GB</b>	<b>1024GB</b>
<b>Total capacity</b>	<b>64TB</b>	<b>1PB</b>
<b>Total server cost</b>	<b>\$4M</b>	<b>\$4M</b>
<b>\$/GB</b>	<b>\$60</b>	<b>\$4</b>


# Motivation: Latency

## Traditional Application



**$\ll 1\mu\text{s}$  latency**

## Web Application




**0.5-10ms latency**

- Large-scale apps struggle with high latency
  - Facebook: can only make 100-150 internal requests per page




# Motivation: Latency

## Traditional Application



**$\ll 1\mu\text{s}$  latency**

## Web Application



~~**$0.5-10\text{ms}$  latency**~~  
 **$5-10\mu\text{s}$**

- RAMCloud goal: large scale **and** low latency
- Enable a new breed of information-intensive applications

# Motivation: Scalability

Relational databases don't scale

Every large-scale Web application has problems:

Facebook: 4000 MySQL instances + 2000 memcached servers

New forms of storage starting to appear:

Bigtable

Dynamo

PNUTS

Sinfonia

H-store

memcached

# Motivation: Technology

Disk access rate not keeping up with capacity:

	Mid-1980's	2009	Change
Disk capacity	30 MB	500 GB	16667x
Max. transfer rate	2 MB/s	100 MB/s	50x
Latency (seek & rotate)	20 ms	10 ms	2x
Capacity/bandwidth (large blocks)	15 s	5000 s	333x
Capacity/bandwidth (1KB blocks)	600 s	58 days	8333x

Disks must become more archival

More information must move to memory

# Why Not a Caching Approach?

Lost performance:

1% misses → 10x performance degradation

Won't save much money:

Already have to keep information in memory

Example: Facebook caches ~75% of data size

Changes disk management issues:

Optimize for reads, vs. writes & recovery

# Why not Flash Memory?

Many candidate technologies besides DRAM

- Flash (NAND, NOR)

- PC RAM

- ...

DRAM enables lowest latency:

- 5-10x faster than flash

Most RAMCloud techniques will apply to other technologies

Ultimately, choose storage technology based on cost, performance, energy, **not volatility**

# Is RAMCloud Capacity Sufficient?

Facebook: 200 TB of (non-image) data today

Amazon:

Revenues/year:	\$16B
Orders/year:	400M? (\$40/order?)
Bytes/order:	1000-10000?
Order data/year:	0.4-4.0 TB?
RAMCloud cost:	\$24-240K?

United Airlines:

Total flights/day:	4000? (30,000 for all airlines in U.S.)
Passenger flights/year:	200M?
Bytes/passenger-flight:	1000-10000?
Order data/year:	0.2-2.0 TB?
RAMCloud cost:	\$13-130K?

Ready today for all online data; media soon

# RAMCloud Overview

Data Durability and Logging

Data Recovery

# RAMCloud Overview

- Data Durability and Logging
- Data Recovery



# Data Durability

## **Want durability with single copy in RAM**

Use disks for replicated backup copies (cheap, non-volatile)

## **Natural to use a logging approach**

Exploit high sequential I/O bandwidth

Avoid high access latencies

## **Scatter backups across cluster**

Fast recovery

High burst write bandwidth

***Next part will discuss recovery from durable storage***

# Data Durability

- **We need durability**


- Servers will fail
- The power will go out
- Failures will be frequent
  - System always in recovery?

- **We need to replicate main memory contents**

- RAM is not feasible
  - Highest performance, but:
    - Assumes we can keep all RAM powered at all times
    - Too expensive: increases cost per bit by replication factor
- Local disks are not feasible
  - Require synchronous writes (latency much too high!)
  - Too slow to recover (single spindle)
  - What if the box dies?

# Buffered Logging

- **Each server maintains a log of updates to its objects**
  - We call the owner server a “*master*”
- **Masters’ log updates are sent to  $R$  backups**
  - RPCs return when backups have updates buffered in RAM
  - Backups batch and write to disk asynchronously
  - Assume for now each master always uses same  $R$  backups
- **Each master is also a backup for other servers**




# Backup Buffer Volatility

- **Problem: Backup buffers are in volatile DRAM**

- Vulnerable until disk flush

- **Solution 1: Synchronous disk writes**

- 2,000 - 8,000 microsecond latency!
- No write batching => very low bandwidth (< 1% of sequential I/O)



- **Solution 2: Synchronous flash SSD writes**

- ~50 - 100 microsecond latency, still a big non-sequential I/O penalty

- **Solution 3: Reduce consistency guarantees**

- “Sorry, we lost your data. Deal with it.”

## **Solution 4: Battery Backups**

Batteries provide enough power to flush buffers

# Log-Structured Backups

- **Problem: Need fast write rates, but have disks**

- RAMCloud is about performance, after all

- **Solution: Log-structure on disks**


- Exploits sequential I/O
- But we need to do cleaning

- **What about log cleaning overheads?**

- All data is in RAM, so no need to re-read for cleaning
  - 50% of the overhead immediately out the window
- Don't need to use disks efficiently
  - Worry about RAM utilization
  - Assume backups have capacity to spare.


# Locating Objects in the Log

- **How do we find objects in the main-memory Log?**
  - Hash table lookup
  - Two cache misses from (tableId, objectId) to object
    - Extremely fast, despite complex memory management scheme



# Locating Objects in the Log

- **How do we find objects in the main-memory Log?**
  - Hash table lookup
  - Two cache misses from (tableId, objectId) to object
    - Extremely fast, despite complex memory management scheme




Hash Table



Log

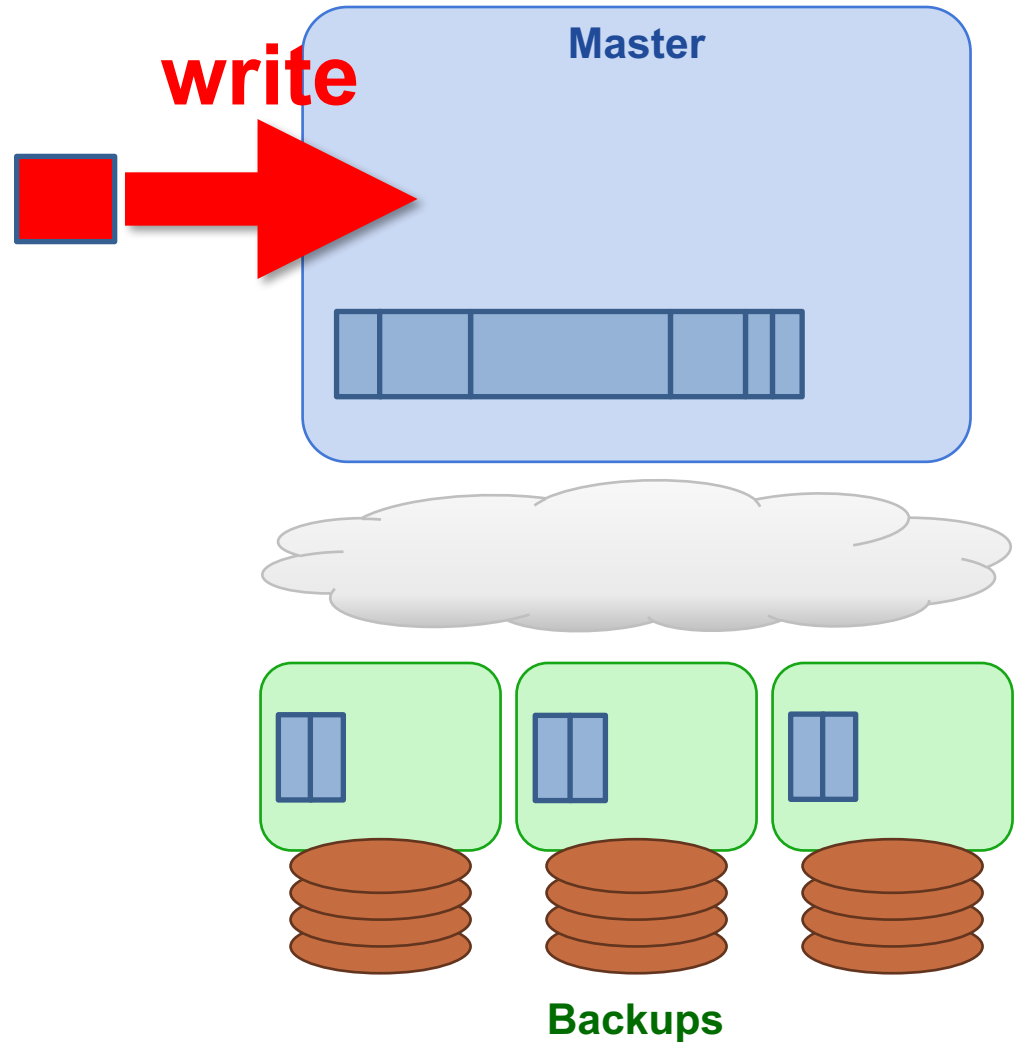
# Locating Objects in the Log

- **How do we find objects in the main-memory Log?**
  - Hash table lookup
  - Two cache misses from (tableId, objectId) to object
    - Extremely fast, despite complex memory management scheme






# Durability with RAM performance




# Durability with RAM performance




# Durability with RAM performance

- Backups buffer update
  - No synchronous disk write




# Durability with RAM performance

- Backups buffer update
  - No synchronous disk write




# Durability with RAM performance

- Backups buffer update
  - No synchronous disk write
- Bulk writes in background
  - Must flush on power loss




# Durability with RAM performance

- Backups buffer update
  - No synchronous disk write
- Bulk writes in background
  - Must flush on power loss
- Pervasive log structure
  - Even RAM is a log
  - Log cleaner



# Durability with RAM performance

- Backups buffer update
  - No synchronous disk write
- Bulk writes in background
  - Must flush on power loss
- Pervasive log structure
  - Even RAM is a log
  - Log cleaner
- Hashtable, key  $\rightarrow$  location



# Data Durability Summary

- Durability with one copy in RAM
- Logging approach for disk I/O utilization
- Backups scattered across cluster for recovery and bursty load



# Outline

- Data Durability and Logging
- Data Recovery

# Data Recovery Overview

- **Master Recovery**
  - 2-Phase
  - Partitioned
- **Failures**
  - Backups
  - Rack/Switch
  - Power
  - Datacenter

# Implications of Single Copy in Memory


- **Problem: Unavailability**
  - If master crashes unavailable until read from disks on backups
  - Read 64 GB from one disk? **10 minutes**
- **Use scale to get low-latency recovery**
  - Lots of disk heads, NICs, CPUs
  - **Our goal:** recover in **1-2 seconds**
    - Is this good enough?
  - Applications just see “hiccups”

# Fast Recovery

- **Problem: Disk bottleneck for recovery**
- **Idea: Leverage many spindles to recover quickly**
  - Data broadly scattered throughout backups
    - Not just great write throughput
    - Take advantage of read throughput
- **Reincarnate masters**
  - With same tables
  - With same indexes
  - Preserves locality


# Fast Recovery: The Problem

- After crash, all backups **read disks** in parallel  
(64 GB/1000 backups @ 100 MB/sec = **0.6 sec, great!**)
- **Collect** all backup data on replacement master  
(64 GB/10Gbit/sec ~ **60 sec: too slow!**)  
Problem: **Network is now the bottleneck!**




## 2-Phase Recovery

- Idea: **Data already in memory** on backups, just need to know **where**




# 2-Phase Recovery

- Phase #1: Recover location info (< 1s)



# 2-Phase Recovery


- **Phase #1: Recover location info (< 1s)**
  - Read all data into memories of backups





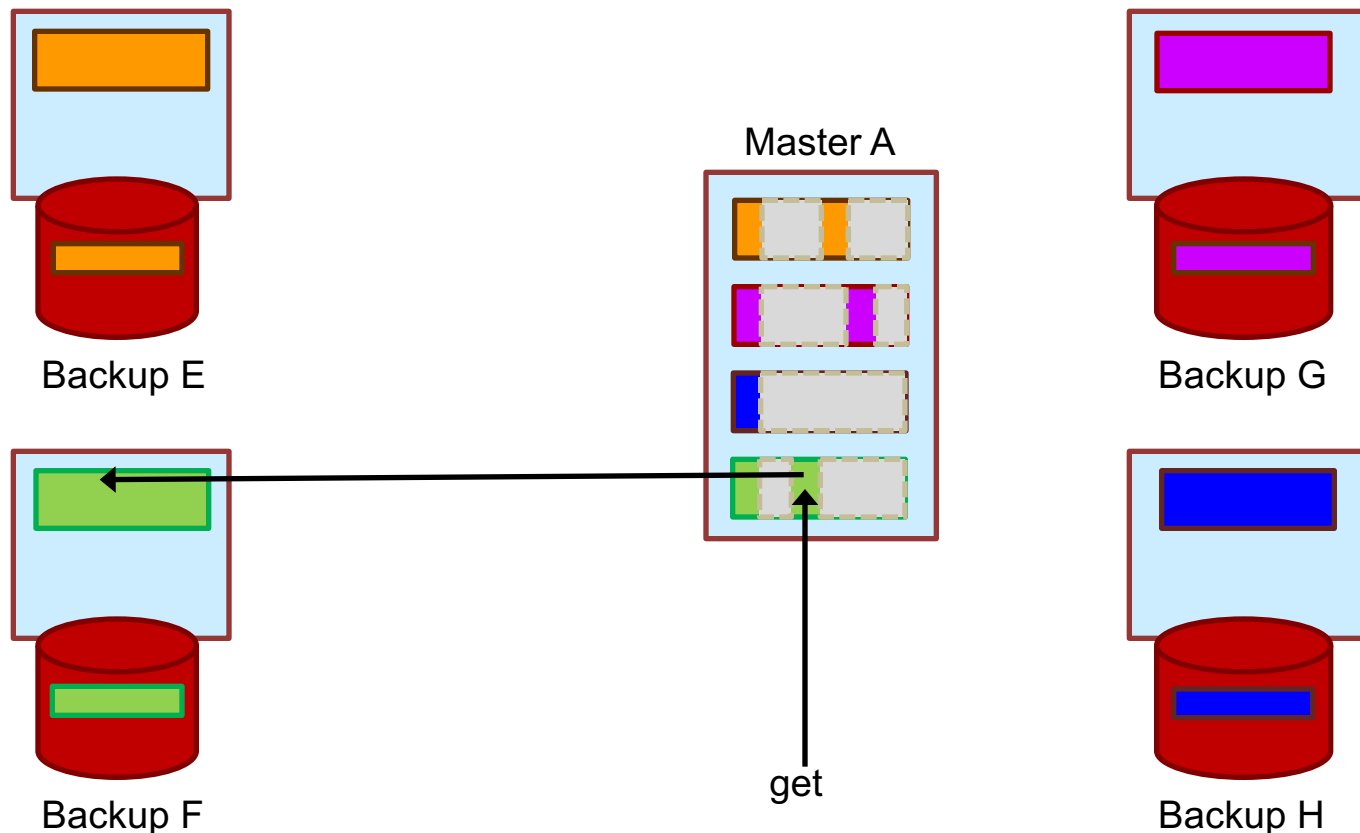
# 2-Phase Recovery

- **Phase #1: Recover location info (< 1s)**
  - Read all data into memories of backups
  - Send **only location info** to replacement master



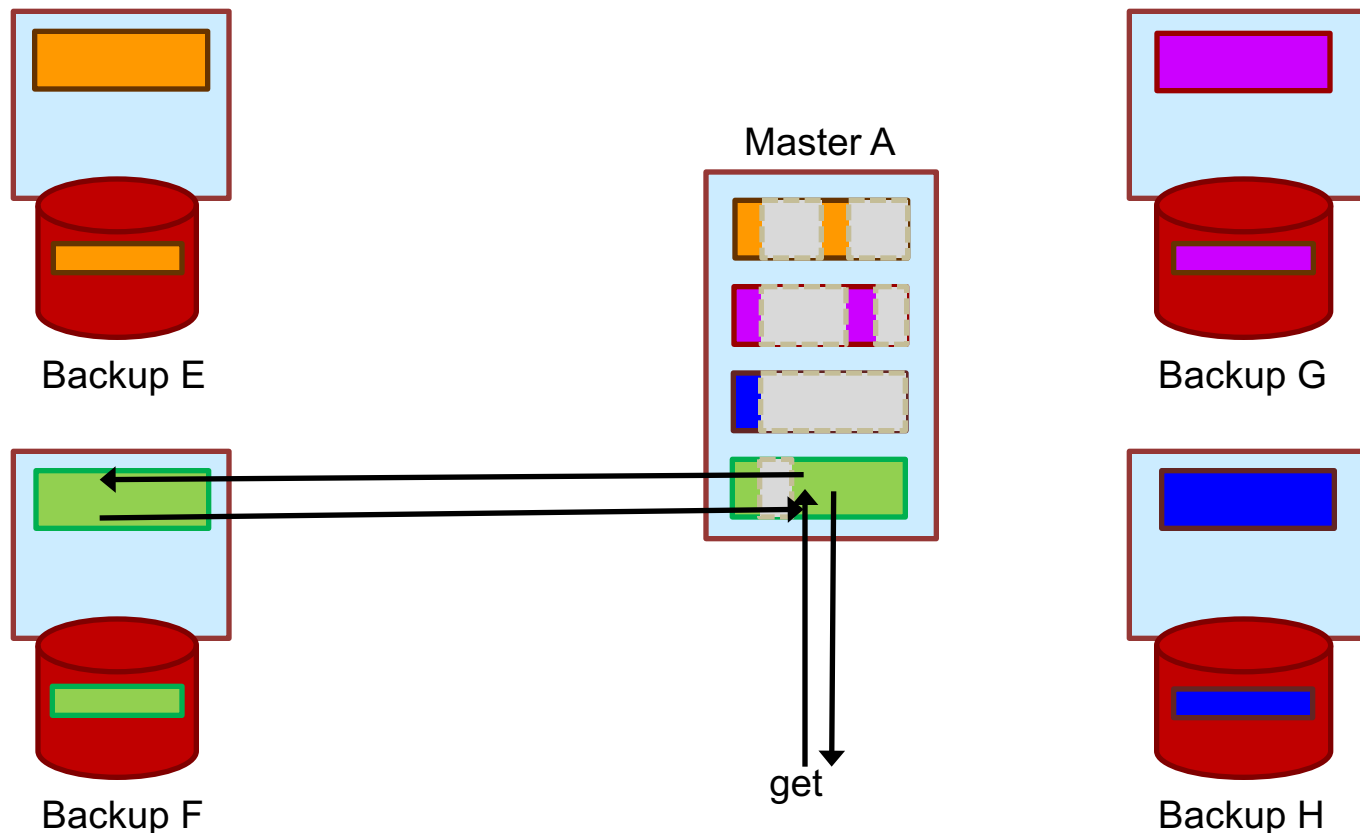
# 2-Phase Recovery

- Phase #2: Proxy & Recover Full Data (~60s)
  - **System resumes operation:**
    - Fetch **on demand** from backups
    - 1 extra round trip on first read of an object
    - Writes are full speed




# 2-Phase Recovery

- Phase #2: Proxy & Recover Full Data (~60s)
  - **System resumes operation:**
    - Fetch **on demand** from backups
    - 1 extra round trip on first read of an object
    - Writes are full speed




## 2-Phase Recovery

- **Phase #2: Proxy & Recover Full Data (~60s)**
  - Transfer data from backups in between servicing requests



# 2-Phase Recovery

- Performance **normal** after **Phase #2** completes



## 2-Phase Recovery: Thoughts

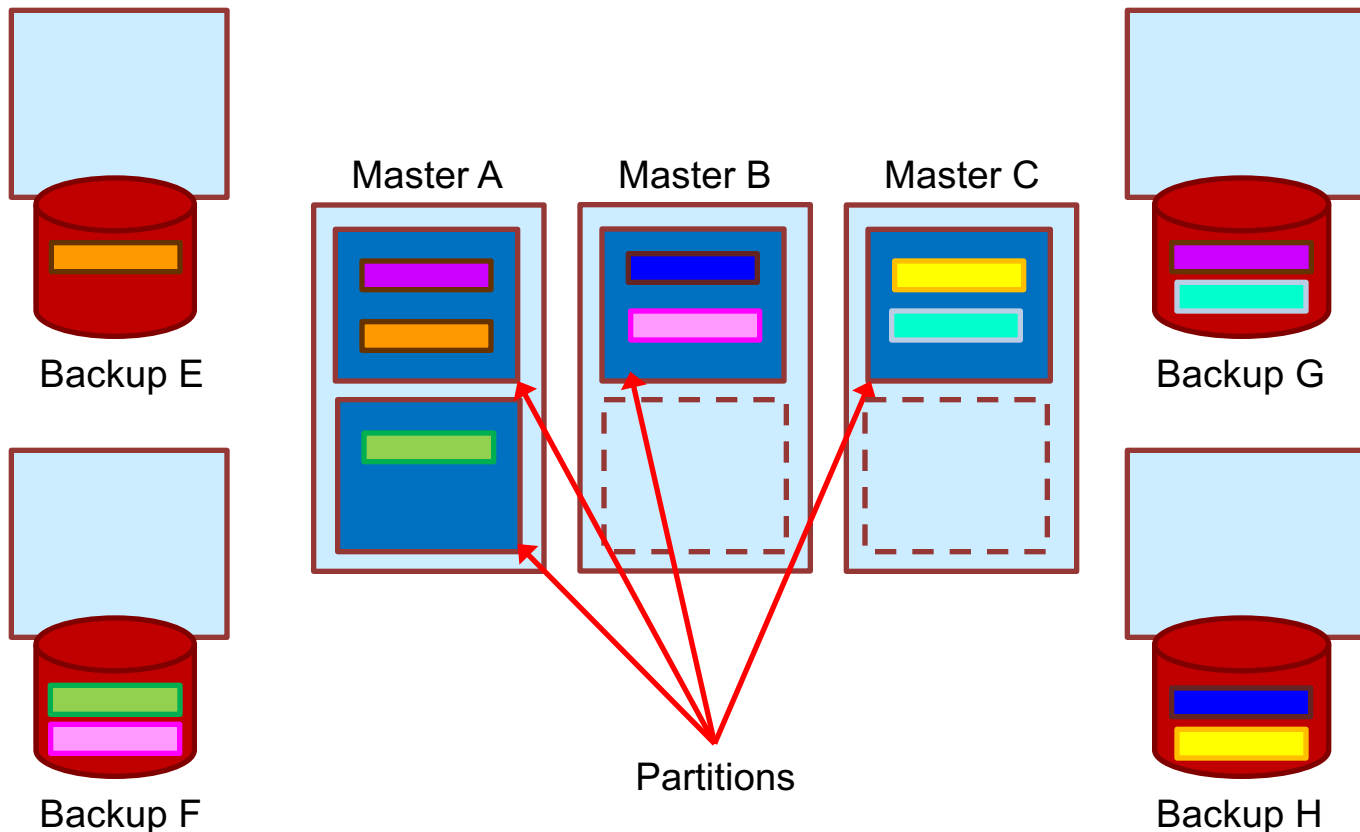
- ✓ **Recovers locality** by recovering machines
- ✗ **Need to talk to **all** hosts**
  - Because backup data for a single master is on all machines
  - How bad is this?
- ✗ **Doesn't deal with heterogeneity**
  - Machine is the unit of recovery
  - Can only recover a machine to one with more capacity
- **Bi-modal Utilization**
  - Must retain pool of empty hosts

## 2-Phase Recovery: Problem

- ✗ **Hashtable inserts become the new bottleneck**
  - Phase #1 must place all objects in the hashtable
  - Master can have 64 million 1 KB objects
  - Hashtable can sustain about 10 million inserts/s
  - 6.4s is over our budget
  - Can use additional cores, but objects could be even smaller
- **Unsure of a way to recover exact master quickly**
  - Constrained by both CPU and NIC
  - Recovery to single host is a bottleneck
- **Another way?**

# Partitioned Recovery


- **Idea: Leverage many hosts to overcome bottleneck**
  - Problem is machines are large so divide them into **partitions**
  - Recover each partition **to a different master**
  - Just like a machine
    - Contains any number of tables, table fragments, indexes, etc.






# Partitioned Recovery

- Load data from disks



# Partitioned Recovery

- Reconstitute **partitions** on many hosts
- 64 GB / 100 partitions = 640 MB
- 640 MB / 10 GBit/s = **0.6s for full recovery**




# Partitioned Recovery: Thoughts

- ✓ **It works: meets availability goals**
  - Can tune time by adjusting partition size
- ✓ **Helps with heterogeneity**
  - Unit of recovery is no longer a machine
- ✗ **Increases host/partition related metadata**
  - Partitions still large enough to provide app locality
- ✗ **Need to talk to **all** hosts**
- **No hot spares**

# Partitioned Recovery: Thoughts

- **Does *not* recover locality**
  - But, no worse than 2-Phase
  - Partitioned approach recovers as fast as Phase #1
  - Can restore locality as fast as Phase #2



# Failures: Backups

- On backup failure the **coordinator broadcasts**
  - More information during coordinator discussion
- All masters **check** their live data
- If objects were on that backup **rewrite** elsewhere (from RAM)
- **No recovery of backups themselves**

# Failures: Racks/Switches

- **Careful selection of backup locations**
  - Write backups for objects **to other racks**
    - As each other
    - As the master
  - Changes as masters recover
    - Can move between racks
  - Masters fix this on recovery
- **Rack failures handled the **same as machine failures****
  - Consider all the machines in the rack dead
- **Question: Minimum RAMCloud that can sustain an entire rack failure and meet recovery goal?**
  - 100 partitions to recover a single machine in 0.6s
  - 50 dead \* 100 partitions, need **5000 machines to make 0.6s**
  - Don't pack storage servers in racks, mix with app servers

# Failures: Power

- **Problem: Objects are buffered temporarily in RAM**
  - Even after the put has returned as successful to the application
- **Solution: All hosts have on-board battery backup**
- **Flush all dirty objects on fluctuation**
  - Any battery should be easily sufficient for this
  - Each master has  $r$  active buffers on backups
  - Buffers are 8MB, for 3 disk copies
    - $3 * 8\text{MB}$  takes 0.24s to flush
- **Question: Cost effective way to get 10-20s of power?**
- **No battery?**
  - Deal with lower consistency
  - Synchronous writes

# Failures: Datacenter

- **Durability guaranteed by disks, no availability**
  - Modulo nuclear attacks
- **No cross-DC replication in version 1**
  - Latency can't be reconciled with consistency
  - Aggregate write bandwidth of 1000 host RAMCloud
    - $100 \text{ MB/s} * 1000 = 1 \text{ Tbit/s}$
- **Application level replication will do much better**
  - Application can batch writes
  - Application understands consistency needs
- **Is this something we need to support?**



# Recovery Summary

- **Use scale in two ways to achieve availability**
  - Scatter reads to overcome disk bottleneck
  - Scatter rebuilding to overcome CPU and network bottlenecks
- **Scale driving lower-latency**
- **Remaining Issue: How do we get information we need for recovery?**
  - Every master recovery involves all backups

# Conclusion

Interesting combination of **scale** and **latency**

Enable more powerful uses of information at scale:

- 1000-10000 clients

- 100TB - 1PB

- 5-10  $\mu$ s latency