

# Read-Log-Update

## A Lightweight Synchronization Mechanism for Concurrent Programming

*Alexander Matveev (MIT) Nir Shavit (MIT and TAU) Pascal  
Felber (UNINE) Patrick Marlier (UNINE)*

Nian Liu  
[nianliu@sjtu.edu.cn](mailto:nianliu@sjtu.edu.cn)

# Concurrent Data Structure

- Multicore environment => Need concurrent data structure
- 2 key concerns for concurrent data structure:

# Concurrent Data Structure

- Multicore environment => Need concurrent data structure
  - 2 key concerns for concurrent data structure:
    - **Unsynchronized** traversals
- 90%** of the time is spent traversing data

# Concurrent Data Structure

- Multicore environment => Need concurrent data structure
- 2 key concerns for concurrent data structure:
  - **Unsynchronized** traversals  
**90%** of the time is spent traversing data
  - **Multi-location** atomic updates  
**Hide** race conditions from programmers

# Read-Copy Update

- Introduced by McKenney
- Linux kernel has 6,500+ RCU calls
  - **Unsynchronized** traversals
  - **Multi-location** atomic updates

# Read-Copy Update

- Introduced by McKenney
- Linux kernel has 6,500+ RCU calls
- **No overhead** for readers
- **Unsynchronized** traversals
- **Multi-location** atomic updates



# Read-Copy Update

- Introduced by McKenney
- Linux kernel has 6,500+ RCU calls
- **No overhead** for readers
- Allow **only single pointer update**
  - **Unsynchronized** traversals
  - **Multi-location** atomic updates






# 2 Key Technique in RCU

- RCU-Duplication
  - Duplicate objects & modify copies
  - Commit through updating a single pointer
- RCU-Epoch
  - Wait for all previous readers to finish
  - Safe to deallocate old objects



# Read-Copy Update


Update C





Lookup C

# Read-Copy Update

Update C




Writer-Lock



Lookup C

# Read-Copy Update


Update C



Lookup C

# Read-Copy Update


Update C



Lookup C

# Read-Copy Update


Update C



Lookup C

# Read-Copy Update


Update C



Lookup C

# Read-Copy Update


Update C



Lookup C

# Read-Copy Update

Update C

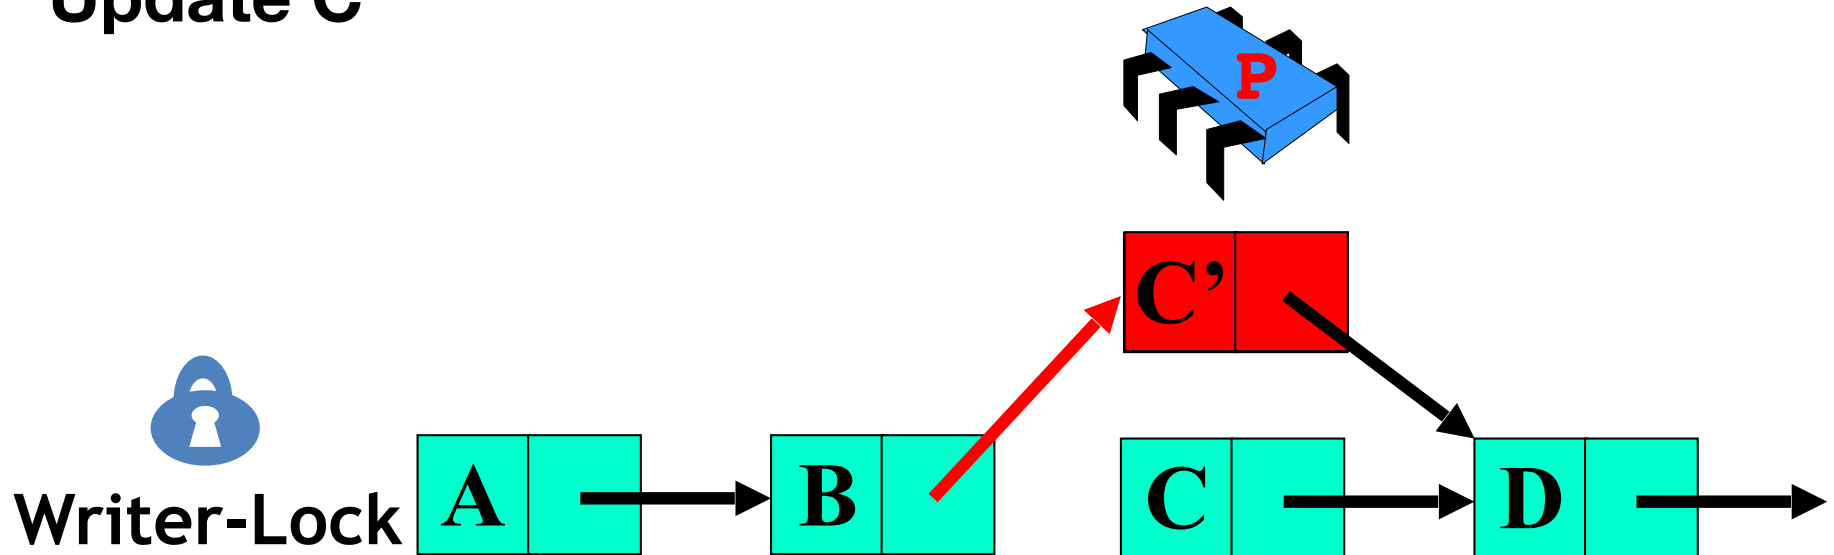


Lookup C




# Read-Copy Update

Update C

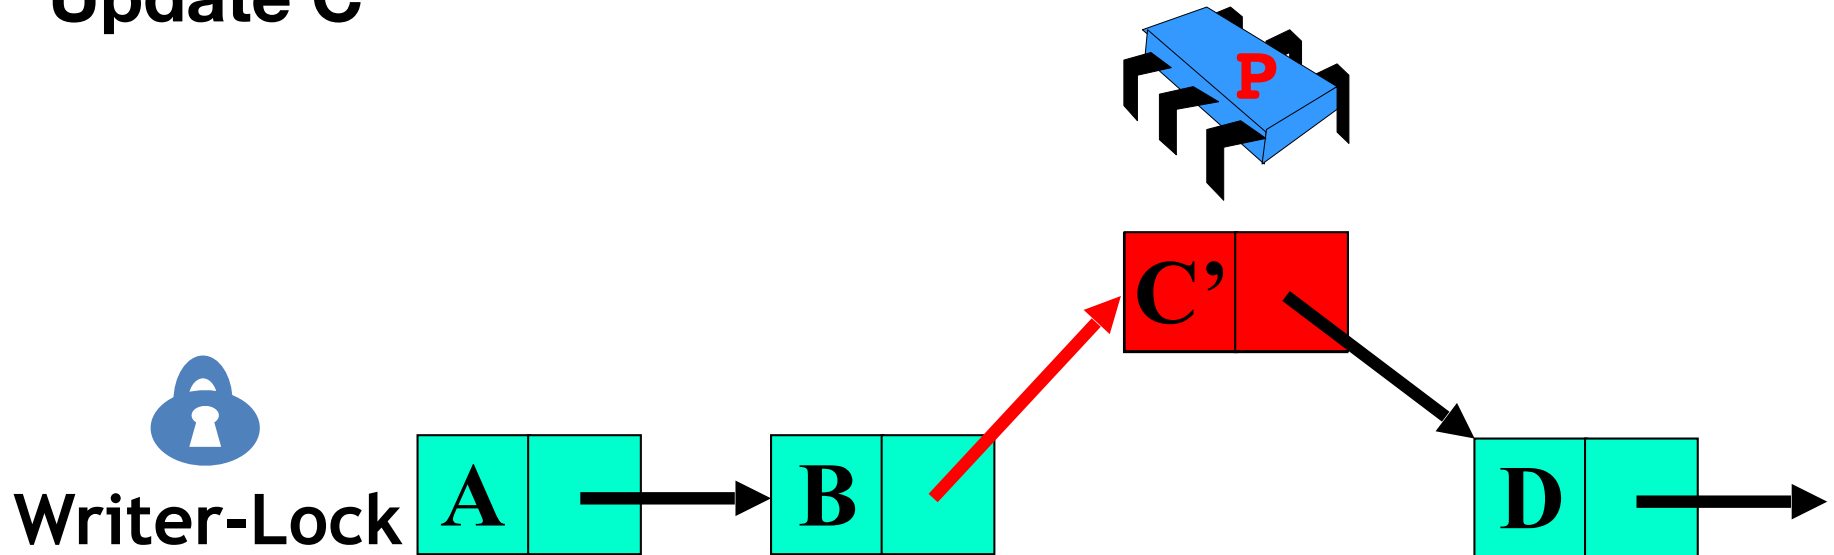


Lookup C




# Read-Copy Update

Update C






Lookup C



# Problems with RCU


Update B & D simultaneous



Lookup B & D

# Problems with RCU


Update B & D simultaneous



Lookup B & D

# Problems with RCU


Update B & D simultaneous



Lookup B & D

# Problems with RCU


Update B & D simultaneous



Lookup B & D

# Problems with RCU


Update B & D simultaneous



Lookup B & D

# Problems with RCU

Update B & D simultaneous




Lookup B & D



# Problems with RCU


Update B & D simultaneous



Lookup B & D


# Problems with RCU

Update B & D simultaneous



Read B' and D

**Inconsistent**



Lookup B & D

# Problems with RCU

- RCU is not capable of dealing with complex data structure

**Double link list / Graph / Resizable hash table**

# Problems with RCU

- RCU is not capable of dealing with complex data structure

**Double link list / Graph / Resizable hash table**

- RCU does not provide concurrency among writers

**A global writer lock** will become a bottleneck

# Problems with RCU

- RCU is not capable of dealing with complex data structure

**Double link list / Graph / Resizable hash table**

- RCU does not provide concurrency among writers

A **global writer lock** will become a bottleneck

- RCU writers have to wait for all readers to complete

Long epoch will **block** the writer

# Introducing RLU

# Introducing RLU

- RCU is not capable of dealing with complex data structure

Update pointer => Commit log by updating **global clock**

- RCU does not provide concurrency among writers
- RCU writers have to wait for all readers to complete

# Introducing RLU

- RCU is not capable of dealing with complex data structure

Update pointer => Commit log by updating **global clock**

- RCU does not provide concurrency among writers

**Fine-grained Lock** for each element

- RCU writers have to wait for all readers to complete



# Introducing RLU

- RCU is not capable of dealing with complex data structure

Update pointer => Commit log by updating **global clock**

- RCU does not provide concurrency among writers

**Fine-grained Lock** for each element

- RCU writers have to wait for all readers to complete

**Deferral** mechanism

# Introducing RLU

- RCU is not capable of dealing with complex data structure

Update pointer => Commit log by updating **global clock**

- RCU does not provide concurrency among writers


**Fine-grained Lock** for each element

- RCU writers have to wait for all readers to complete

**Deferral** mechanism

# Read Log Update

Update B & D simultaneous






# Read Log Update

Update B & D simultaneous

Writer  
Timestamp =  $\infty$

Global  
Timestamp = 22




Lookup B & D

# Read Log Update

Update B & D simultaneous

Writer  
Timestamp =  $\infty$

Global  
Timestamp = 22



Lookup B & D


Local Timestamp = 22

# Read Log Update

Update B & D simultaneous

Writer  
Timestamp =  $\infty$

Global  
Timestamp = 22



Lookup B & D


Local Timestamp = 22

# Read Log Update

Update B & D simultaneous

Writer  
Timestamp =  $\infty$

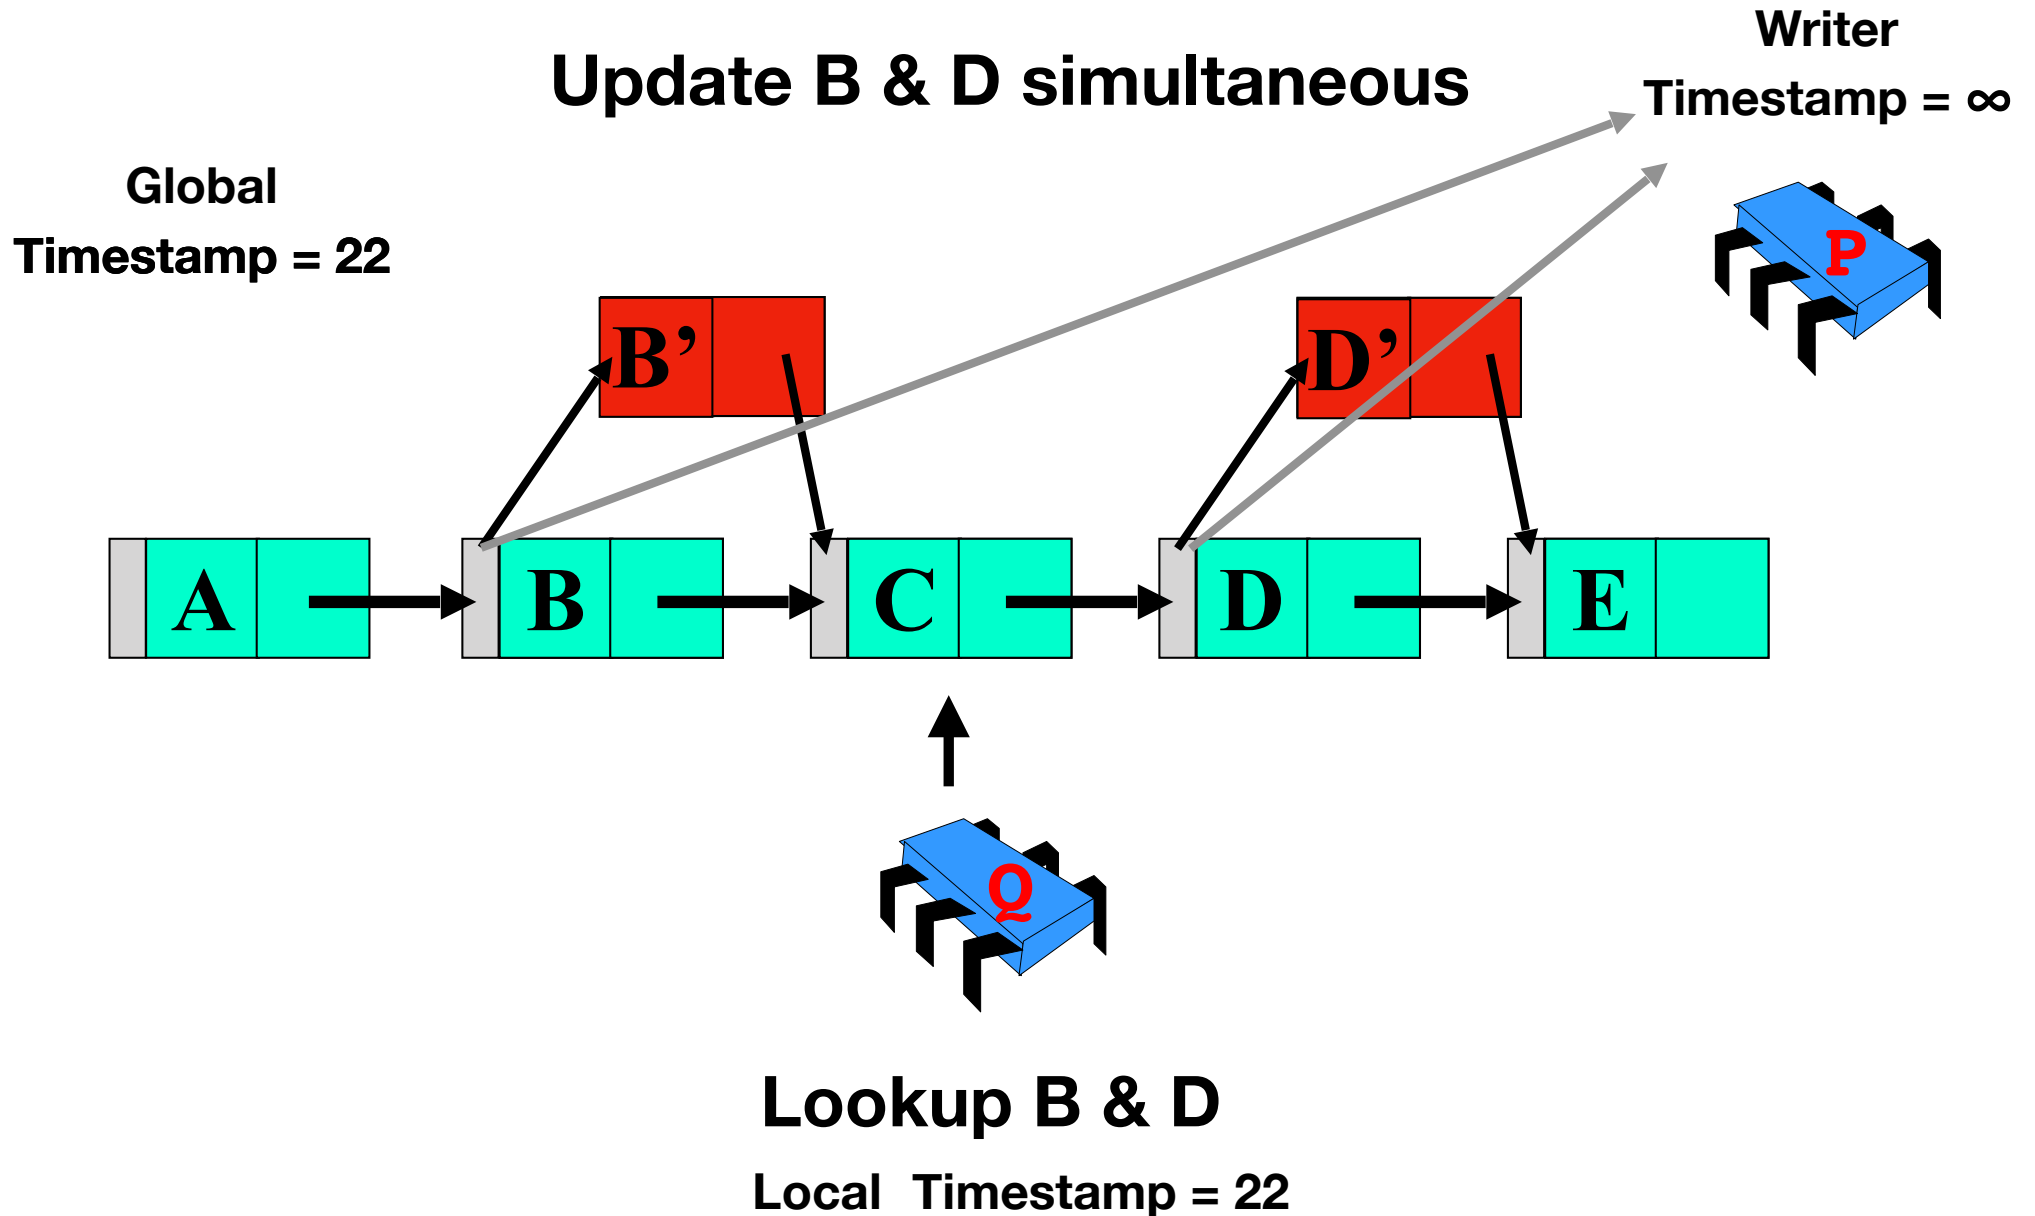
Global  
Timestamp = 22



Lookup B & D

Local Timestamp = 22

# Read Log Update





# Read Log Update


Update B & D simultaneous

Writer

Timestamp = 22 + 1

Global

Timestamp = 22 + 1



Lookup B & D

Local Timestamp = 22

# Read Log Update


Update B & D simultaneous

Writer

Timestamp = 23

Global

Timestamp = 23



Lookup B & D

Local Timestamp = 22



# Read Log Update

Update B & D simultaneous

Writer

Timestamp = 23

Global  
Timestamp = 23



Anyone update D ?




Equal or smaller timestamp?



Lookup B & D

Local Timestamp = 22



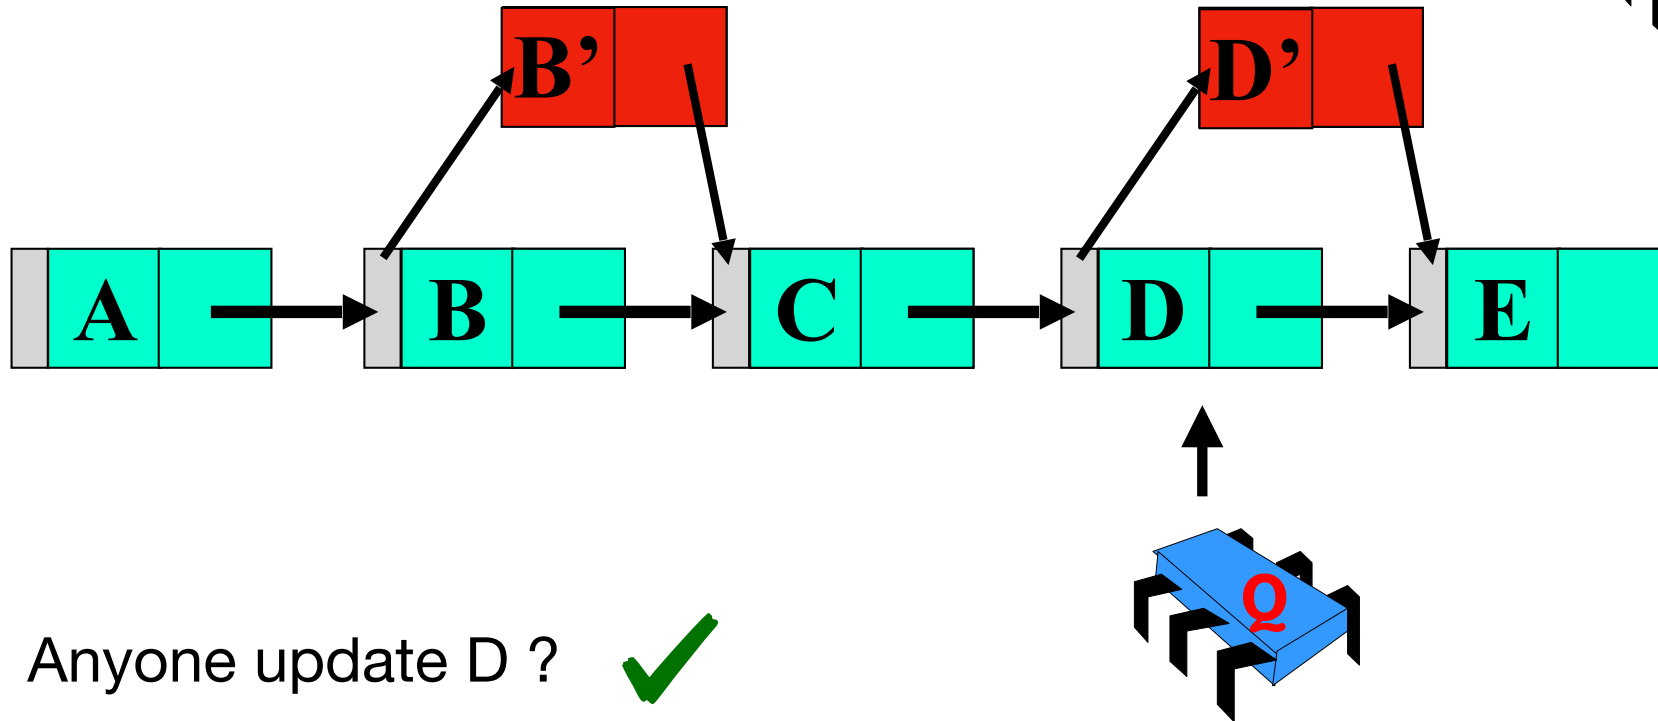
# Read Log Update

Update B & D simultaneous

Writer

Timestamp = 23

Global  
Timestamp = 23



Anyone update D ?



Equal or smaller timestamp?



Lookup B & D

Local Timestamp = 22



# Read Log Update

Update B & D simultaneous

Writer

Timestamp = 23

Global  
Timestamp = 23






# Read Log Update

Update B & D simultaneous

Writer

Timestamp = 23

Global  
Timestamp = 23



Lookup B & D

Timestamp = 23

Anyone update B ?



Equal or smaller timestamp?






# Read Log Update

Update B & D simultaneous

Writer

Timestamp = 23

Global  
Timestamp = 23



Lookup B & D

Timestamp = 23

Anyone update B ?



Equal or smaller timestamp?




# Read Log Update

Update B & D simultaneous

Writer

Timestamp = 23

Global  
Timestamp = 23



Anyone update D ?



Equal or smaller timestamp?



Lookup B & D

Timestamp = 23




# Read Log Update

Update B & D simultaneous

Writer

Timestamp = 23

Global  
Timestamp = 23



Wait until all previous readers finish their jobs

Readers' timestamp < Current writers'



# Read Log Update

Update B & D simultaneous

Writer

Timestamp = 23

Global  
Timestamp = 23





Write back logs when previous readers finish their jobs

# Read Log Update

Update B & D simultaneous

Timestamp =  $\infty$

Global  
Timestamp = 23





Reset write timestamp

# Read Log Update

Update B & D simultaneous

Global  
Timestamp = 23



Swap to another log buffer

# Most exciting part in RLU

- 2 key concerns for concurrent data structure:

- **Unsynchronized** traversals



- **Multi-location** atomic updates



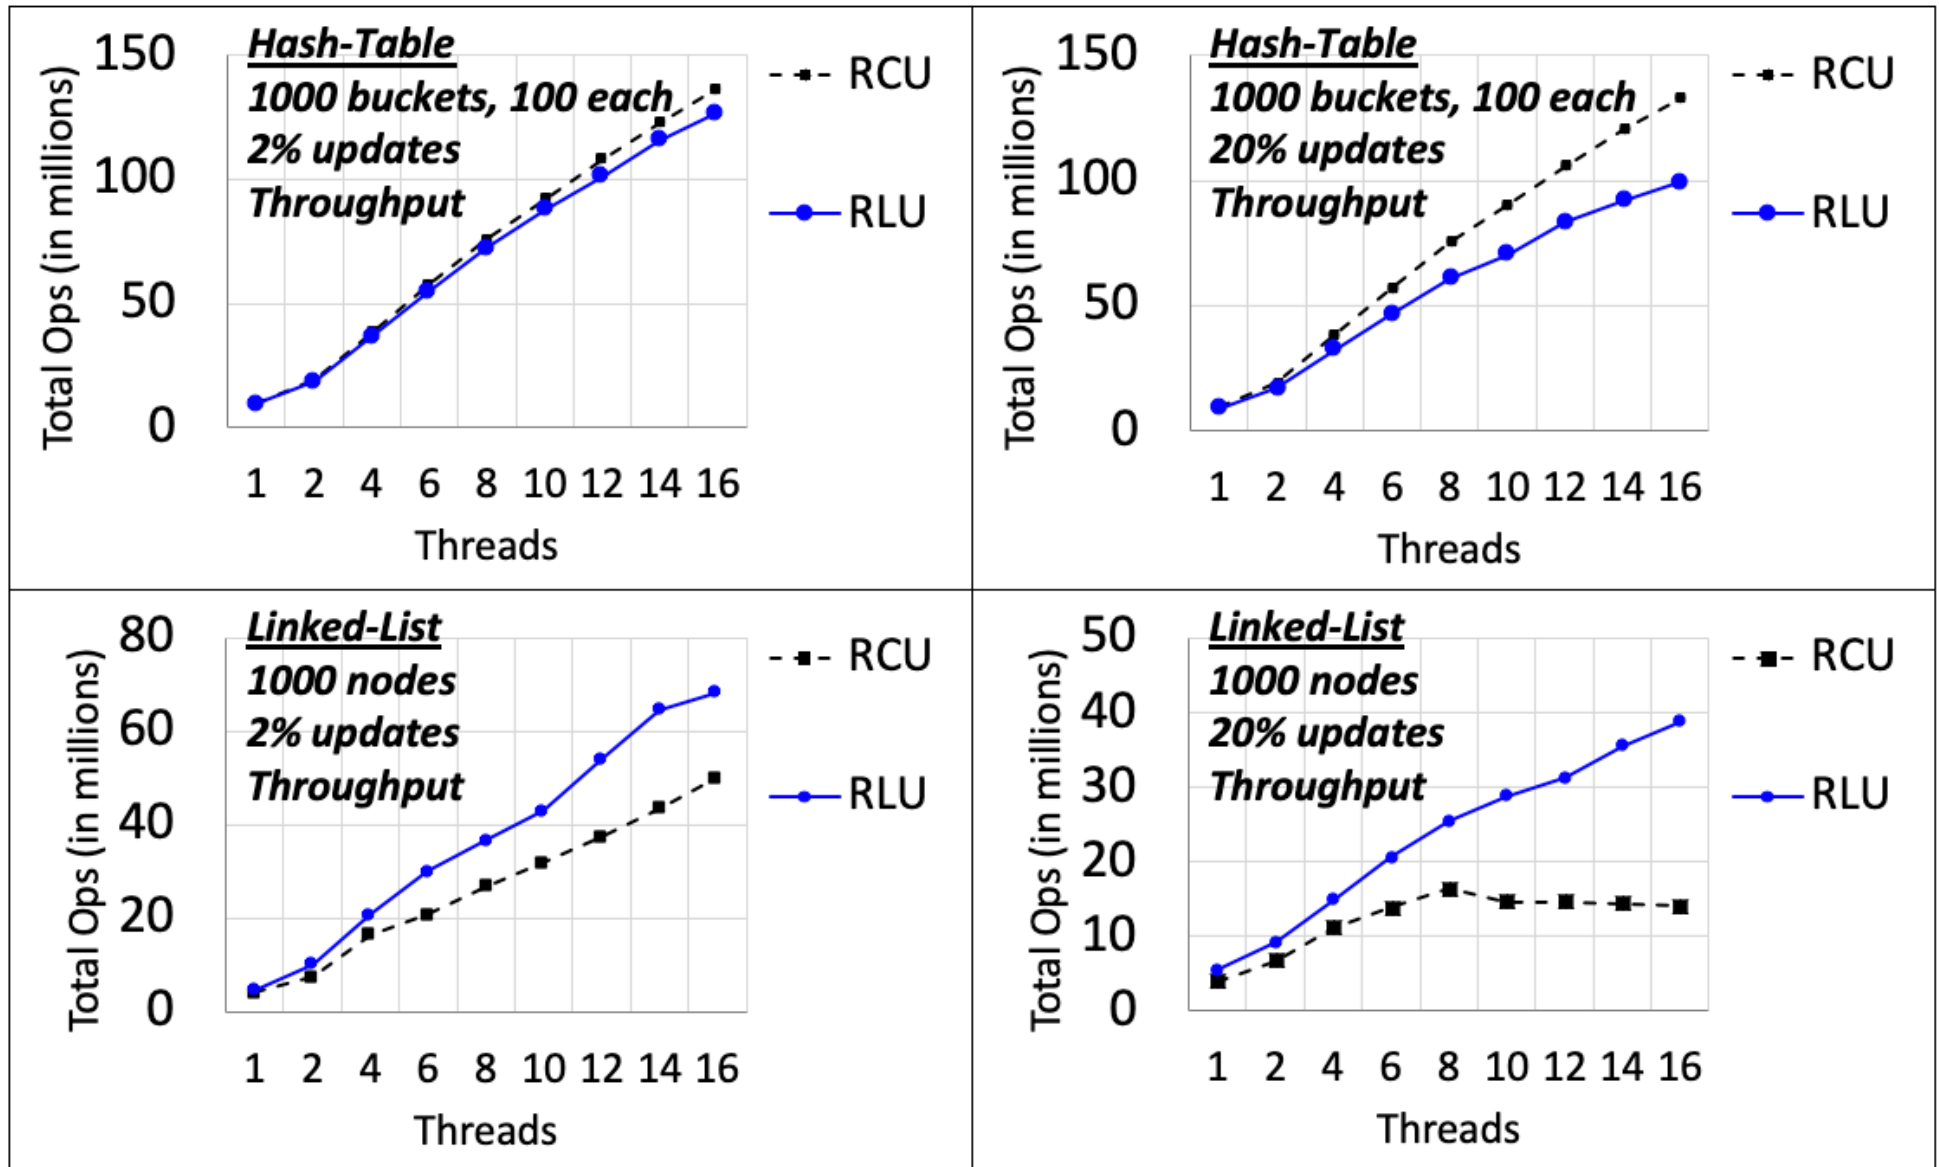
- Easier to be used in complicated applications

- Double link list


- Kyoto Cabinet

# Evaluation

# Microbenchmarks



# Kyoto Cabinet





# Conclusion

- Single pointer update => Global clock update
- Stronger semantic compared to RCU
- Significantly lower intellectual cost to the programmer
- No extra overhead on reader's critical path
- Achieve similar (even better) performance with RCU



**Thanks**  
**Q & A**