

# Network Function Virtualization

Haibo Chen

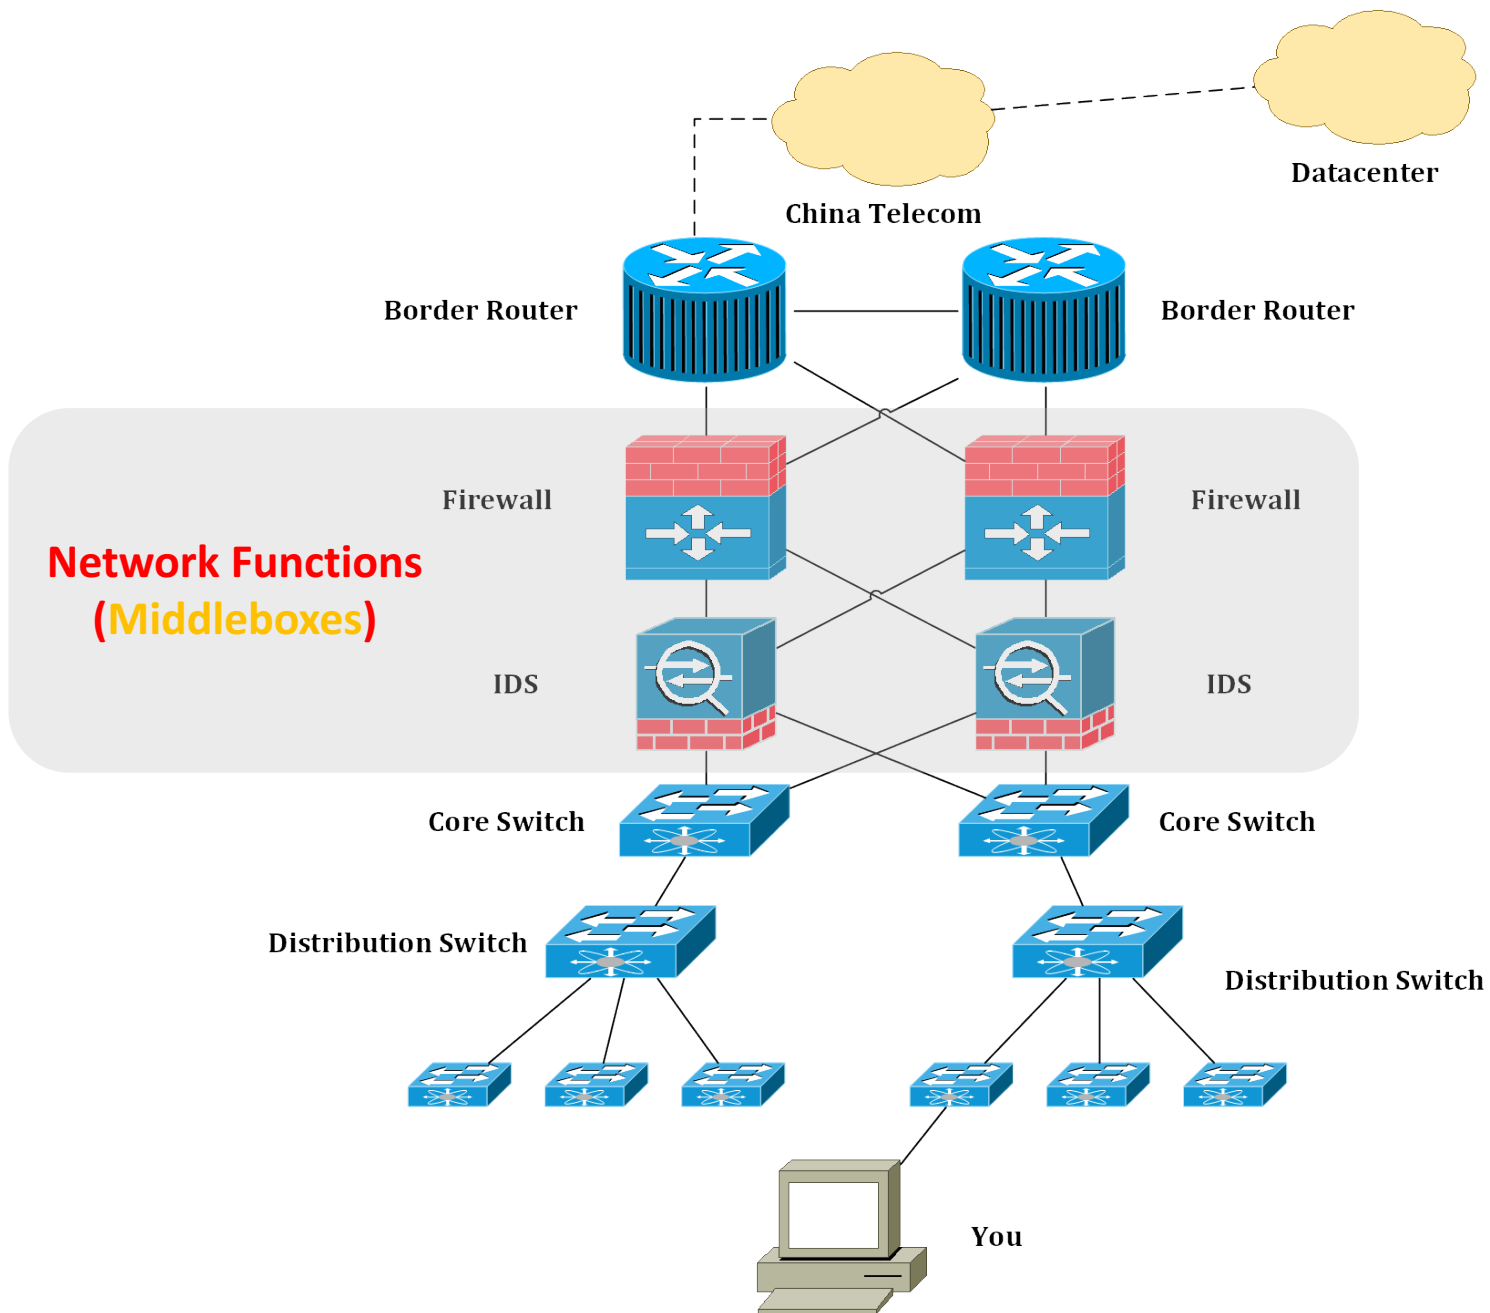
Institute of Parallel and Distributed Systems (IPADS)


Shanghai Jiao Tong University

[http://ipads.se.sjtu.edu.cn/haibo\\_chen](http://ipads.se.sjtu.edu.cn/haibo_chen)

# How do you access servers?








WAN Optimizer




**\$235,000**



Firewall




**\$4,923**



Intrusion Detection  
System




**\$43,700**



WAN Optimizer



\$235,000



Firewall



\$4,923



Intrusion Detection System



\$43,700



Intel E5-2650 v3 with 20 cores, 40 Gbps Ethernet NIC, 64 GB RAM

¥45,000

# Goal of NFV

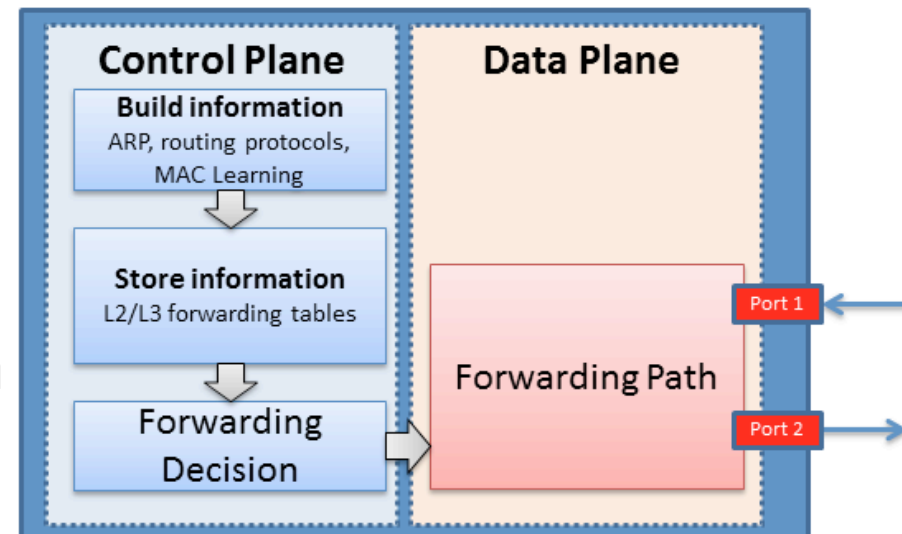
- Save money
  - Commodity server hardware
  - Workload consolidation
  - Lower power consumption
  - Simplified network maintenance
- Make money
  - Accelerated service deployment
  - Network infrastructure as a service

# What is NFV?

- Network Function
  - Middleboxes (firewall, IDS, WAN optimizer)
- Virtualization
  - Commodity hardware
  - Consolidation
- **Network Services** running on **Cloud Infrastructure**

# Relationship with SDN

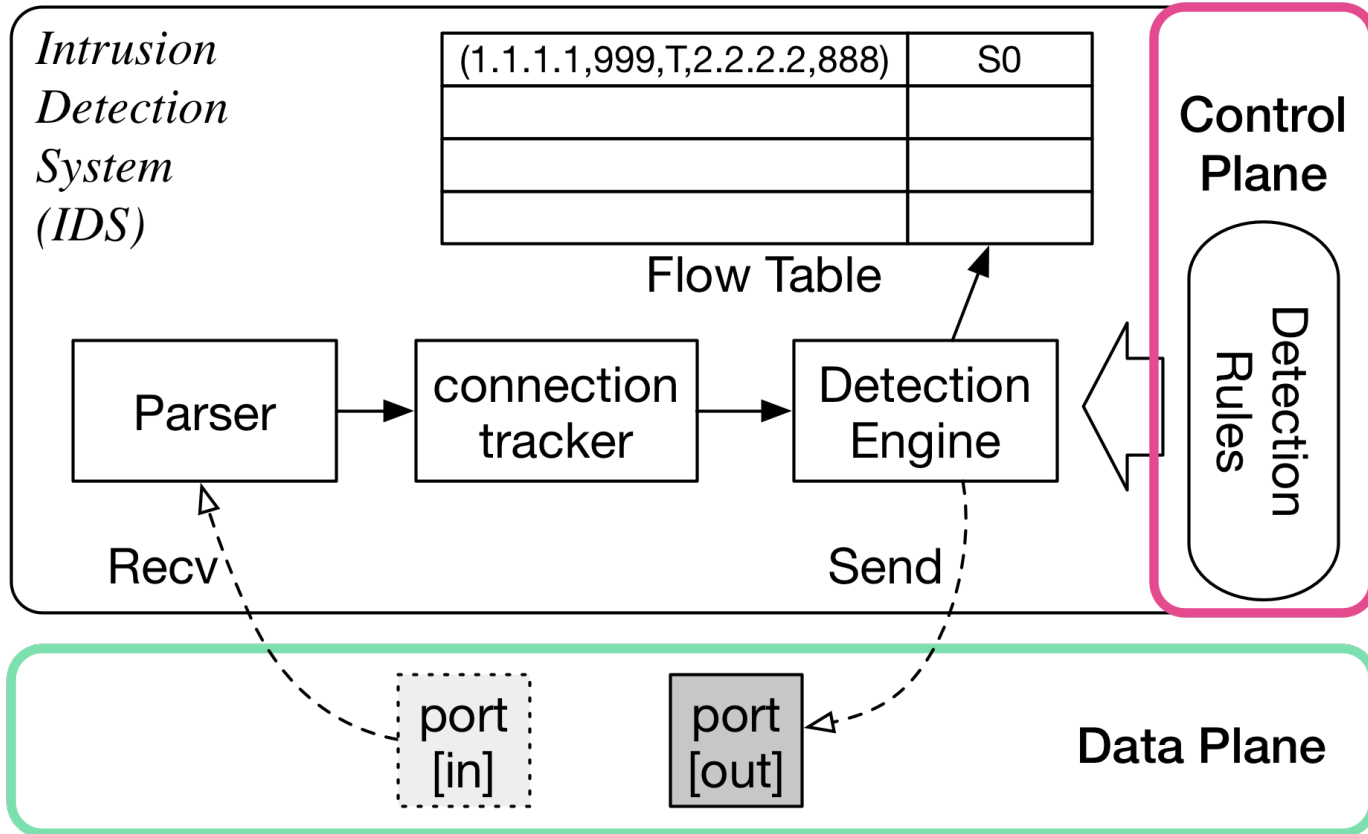
- SDN: Software-defined networking
  - Focus on control plane innovation/virtualization
- NFV: network function virtualization
  - Focus on virtualizing data plane
- SDN & NFV
  - Can innovate separately
  - Or together
  - Like policy vs. mechanism






# NFV: CHALLENGES

# Challenges: Data Plane Performance



# Challenges: Data Plane Performance



# Challenges: Consolidation

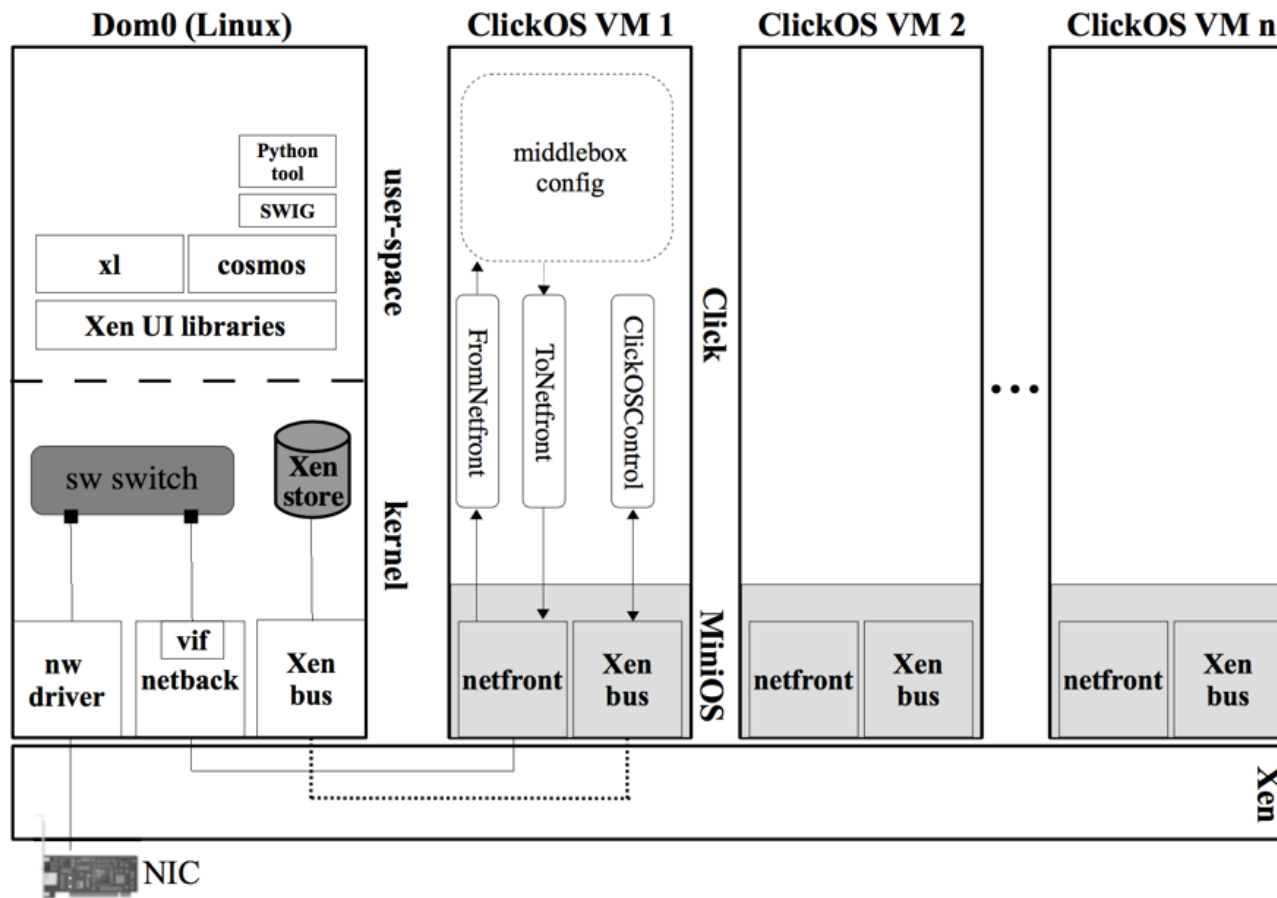
Middlebox	Functionality	Examples
NIDS	Monitor network activities for malicious behaviors or policy violations and report to operations staff	Bro IDS, Snort
Firewall	A firewall based on packet filter drops packets matching specific patterns and forwards the rest to next middlebox	iptables, pf
Load Balancer	Load balancers distribute network flows across several down- stream middleboxes	HAProxy, LVS
NAT	Remap one IP address space to another	iptables

# Challenges: Consolidation

- Process
  - No performance penalty
  - No performance isolation
- Virtual machine
  - Virtualization overhead
  - Performance isolation
  - Easy deployment
  - Mature technology
- Container
  - No performance penalty
  - Performance isolation
  - Easy deployment

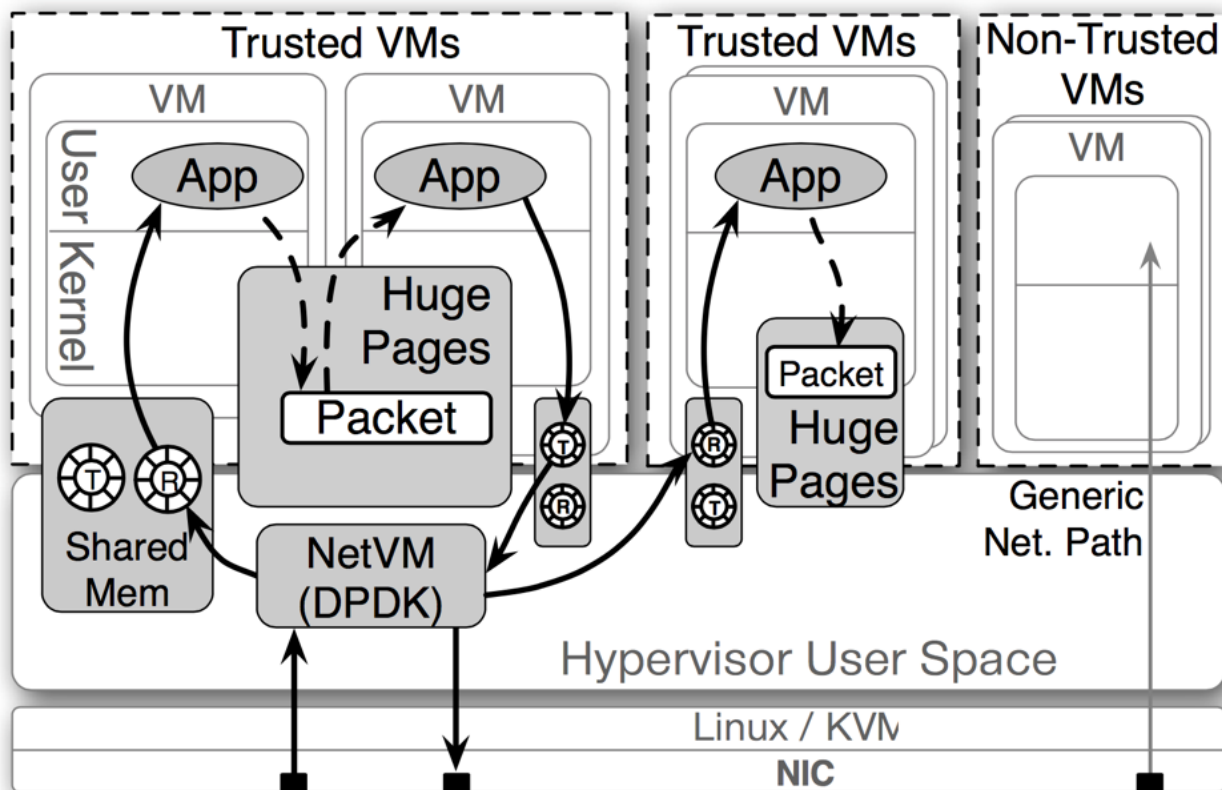
# Challenges: Consolidation

- Virtual machine ClickOS (NSDI'14)



# Challenges: Consolidation

- Virtual machine **NetVM (NSDI'14)**

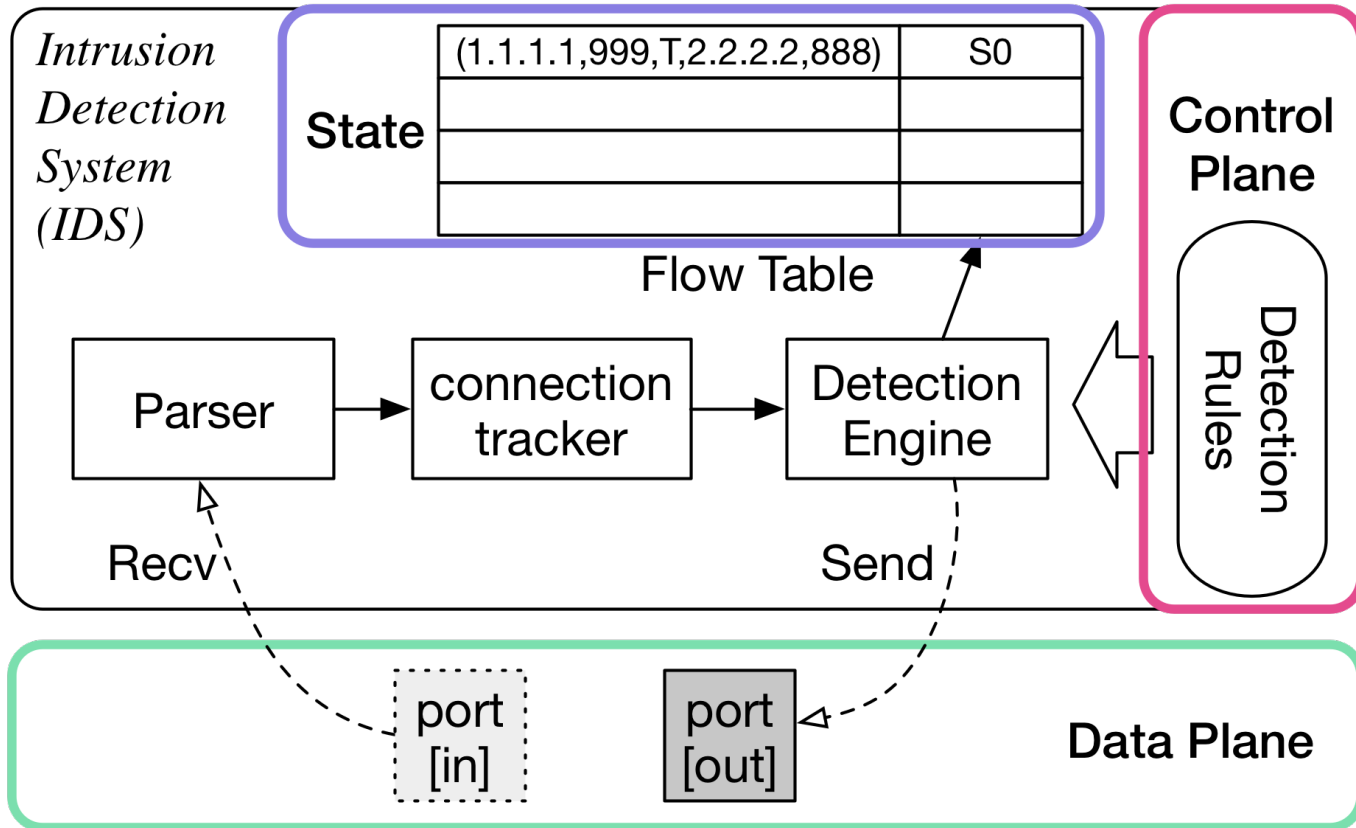


# Challenges: Scalability

- Dynamic network
  - Workload
  - Reconfiguration
  - Resource allocation
- Scalability
  - Scale up
    - Multicore hardware
    - NUMA aware
  - Scale out
    - Replication
    - Migration



# Challenges: State Management



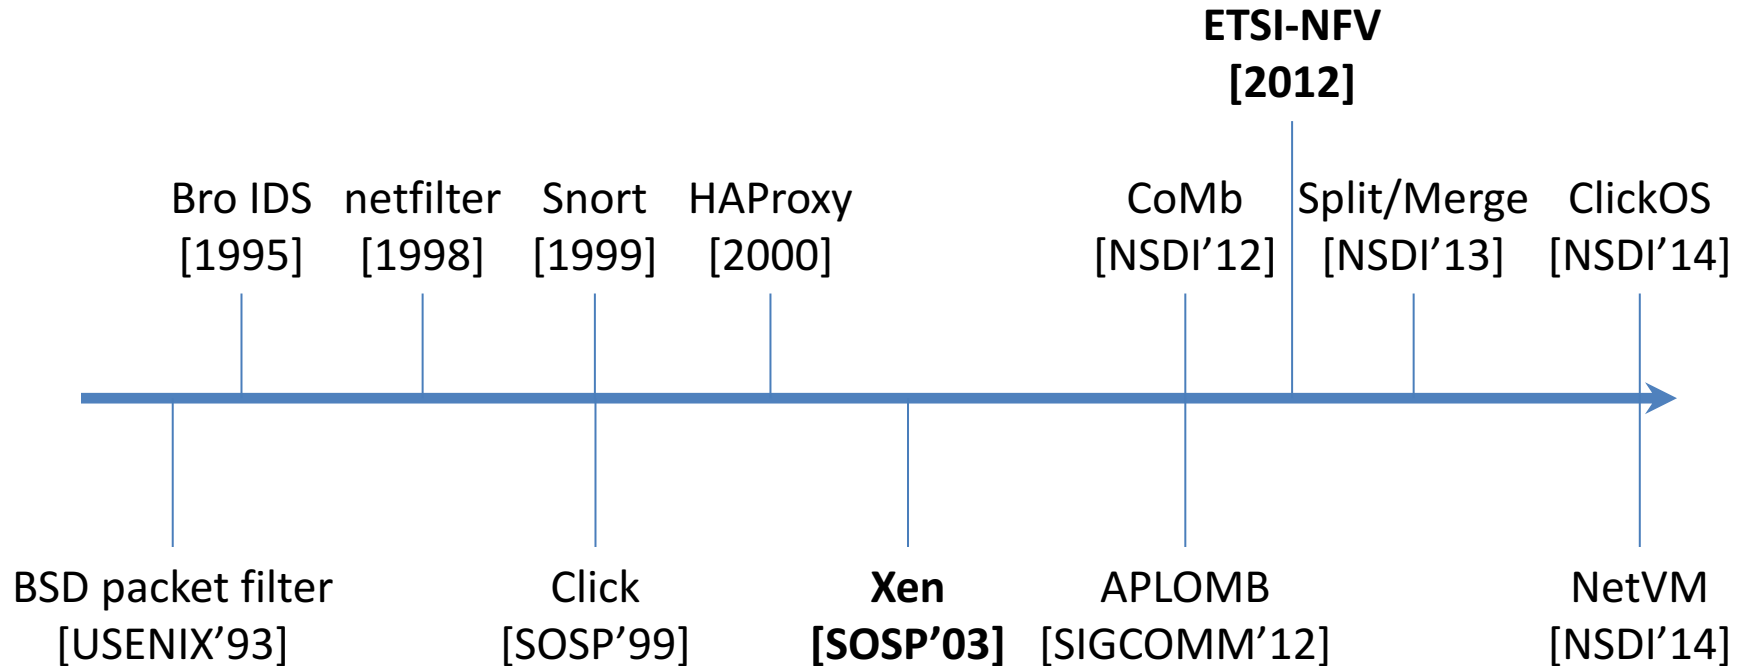
# Challenges: State Management

- State consistency
  - Distributed instances
  - Recovery and migration
- Requirements
  - Classified states: internal & external
  - Transactional boundary
  - Move states between replicas
  - Correctly route traffic

# NFV: History

- Software middlebox
  - IDS: Bro, Snort
  - Firewall: iptables (1998), BSD pf (USENIX'93)
  - Load Balancer: HAProxy, LVS
  - NAT: iptables
- Click (SOSP'99)
- Middlebox consolidation
  - NIDS Cluster (RAID'07)
  - APLOMB (SIGCOMM'12)
  - CoMb (NSDI'12)
  - ClickOS (NSDI'14)

# NFV: History



Software Packet Processing

Network Function Virtualization

# NFV: SUMMARY

# NFV: Summary



at&t  
arrr




HUAWEI



CISCO™



Alcatel•Lucent



JUNIPER®  
NETWORKS

# NFV: Summary

- Virtualization of network appliances
- Performance and flexibility
- System design
  - Multicore and NUMA hardware
  - Distributed processing
  - I/O optimization

# **The Click Modular Router**

Robert Morris, Eddie Kohler, John  
Jannotti, M. Frans Kaashoek



# Background

- New network functions
  - NIDS
  - Random early detection dropping policy
  - Deep packet inspection
  - GFW
- Commercial off-the-shelf (COTS) hardware
  - Expensive
  - Closed
  - Static
  - Inflexible

# Motivation

- Flexible
  - New functionality
- Modular
  - Combine elements
- Open
  - Allow developers to add elements
- Efficient
  - Comparable with hardware


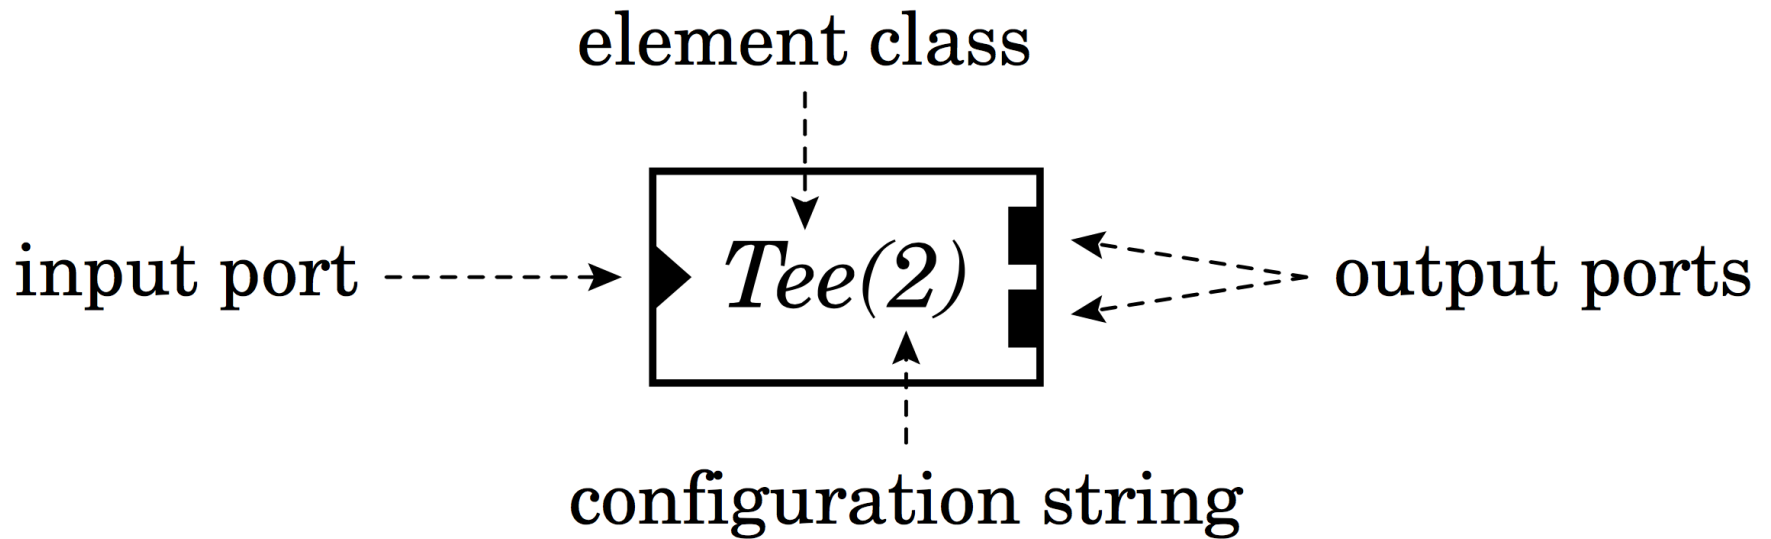
# Click: Overview

- A modular design to allow easy construction of routers and other middleboxes
  - Elements
    - Push
    - Pull
  - A domain-specific language (DSL)
    - Declare elements
    - Connect elements
  - Lightweight framework
    - Orchestrate elements into a router

# Elements

- Building block of Click router
  - C++ classes
  - Scheduled by framework
- Ports
  - Endpoints of connections
  - Associated with push or pull function (or *agnostic*)
- Configuration
  - Parameters to initialize elements
  - Elements specific

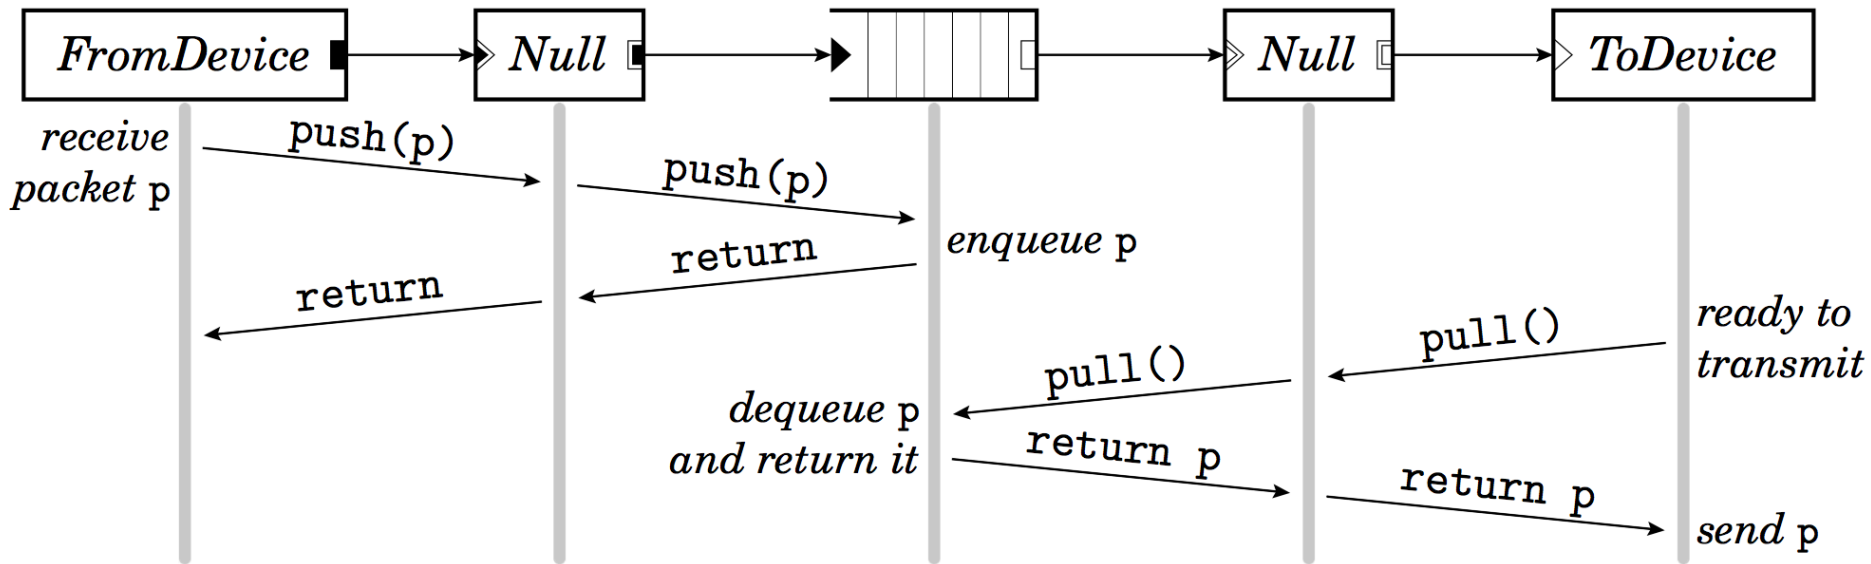
# Element Example



# C++ Class of Elements

```
class NullElement : public Element {  
    public:  
    NullElement()  
        { add_input(); add_output(); }  
    const char *class_name() const  
        { return "Null"; }  
    PushOrPull default_processing() const  
        { return AGNOSTIC; }  
    NullElement *clone() const  
        { return new NullElement; }  
    void push(int port_number, Packet *p)  
        { output(0).push(p); }  
    Packet *pull(int port_number)  
        { return input(0).pull(); }  
};
```

# Push and Pull




## Push


From source element to downstream  
Triggered by event, e.g. packet arrival  
Denoted by solid square or triangle

## Pull

From destination element to upstream  
Packet transmission  
Denoted by empty square or triangle



# More Examples





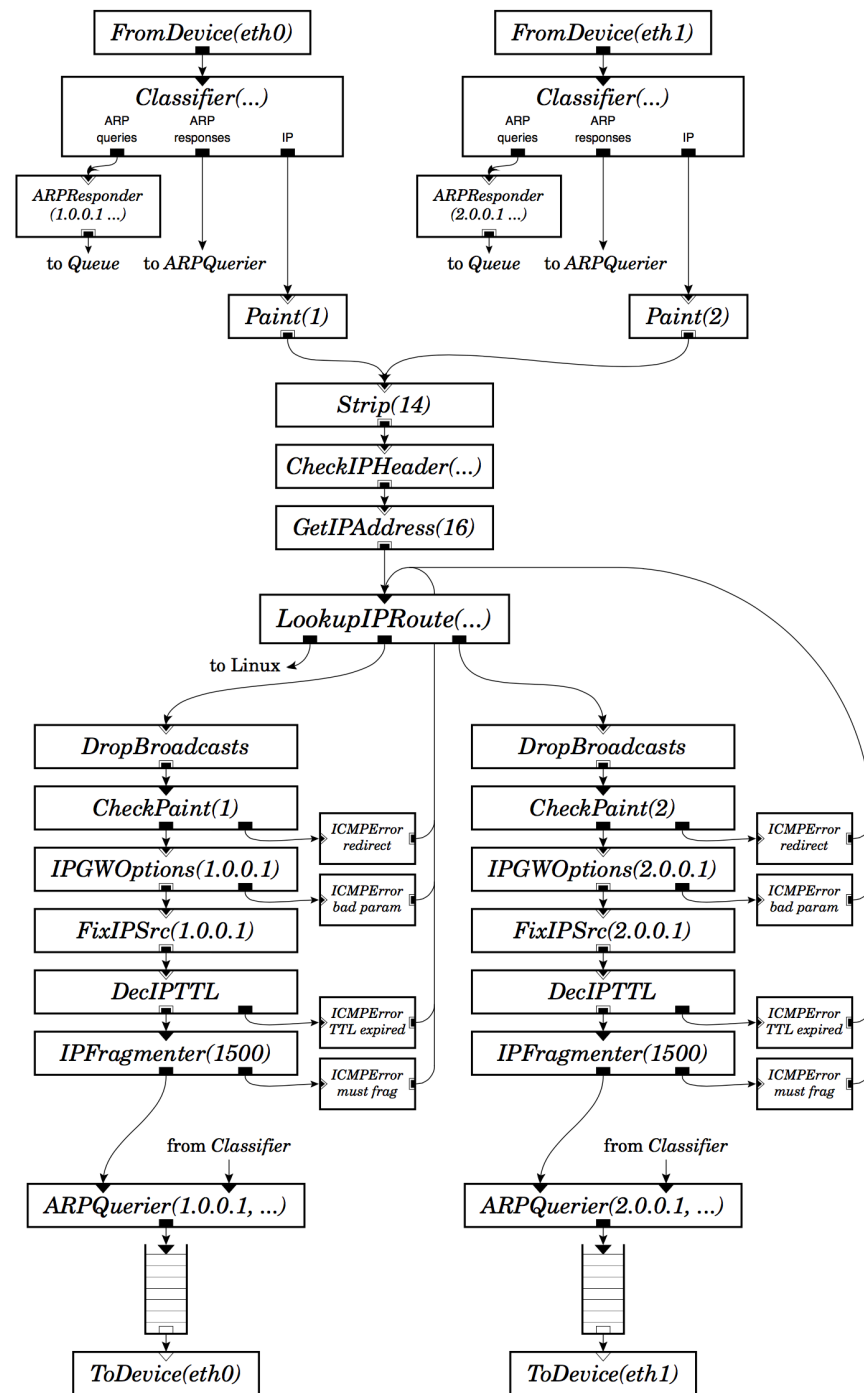
# Packet Counter



```
// Declare three elements ...  
src :: FromDevice(eth0);  
ctr :: Counter;  
sink :: Discard;  
// ...and connect them together  
src -> ctr;  
ctr -> sink;
```

# Ping Responder

```
1 define($IP 192.168.13.235);
2 define($MAC 00-15-17-15-5d-75);
3
4 tap :: KernelTap(192.168.13.235/16, ETHER $MAC);
5
6 c :: Classifier(
7     12/0806 20/0001, // ARP Requests goes to output 0
8     12/0806 20/0002, // ARP Replies to output 1
9     12/0800, // ICMP Requests to output 2
10    -);
11
12 arpq :: ARPQuerier($IP, $MAC);
13 arpr :: ARPResponder($IP $MAC);
14
15 tap-> c;
16 c[0] -> ARPPrint -> arpr -> Print("arpr", 60) -> tap;
17 c[1] -> [1]arpq;
18 Idle -> [0]arpq;
19 arpq -> Print("arpq", 60) -> ARPPrint -> tap;
20 c[2] -> CheckIPHeader(14) -> ICMPPingResponder() -> EtherMirror() -> tap;
21 c[3] -> Discard;
```



# Scheduling


- Single thread in the original paper
- Only some elements need to be scheduled
- Click schedules the CPU with a task queue
  - Flexible and lightweight stride scheduling [Waldspurger'95]
  - Element is the unit of CPU scheduling
  - Most elements are never placed on the task queue
  - Elements frequently initiates push or pull should be scheduled (`FromDevice`, `ToDevice`)

# Implementation

- Userlevel
  - Run as normal thread
  - FromDevice: TUN/TAP, libpcap
- Kernel module
  - Run as kernel thread
  - Loops over the task queue and run each task
  - Only interrupts can preempt this thread
  - Voluntarily gives up CPU to Linux
- ClickOS [NSDI'14]
  - Run in MiniOS on Xen
  - Use netmap [ATC'12] to get packets

# Maximum Loss-Free Forwarding Rate

- Click IP router
  - 8 interfaces
  - 161 elements
  - 64 bytes packets size
- Forwarding rate versus input rate
- MLFFR reflects the router's behavior under load



# Adding Network Functions

## Simple

Passes packets from one interface to another

## Switch

L2 switch

## IP


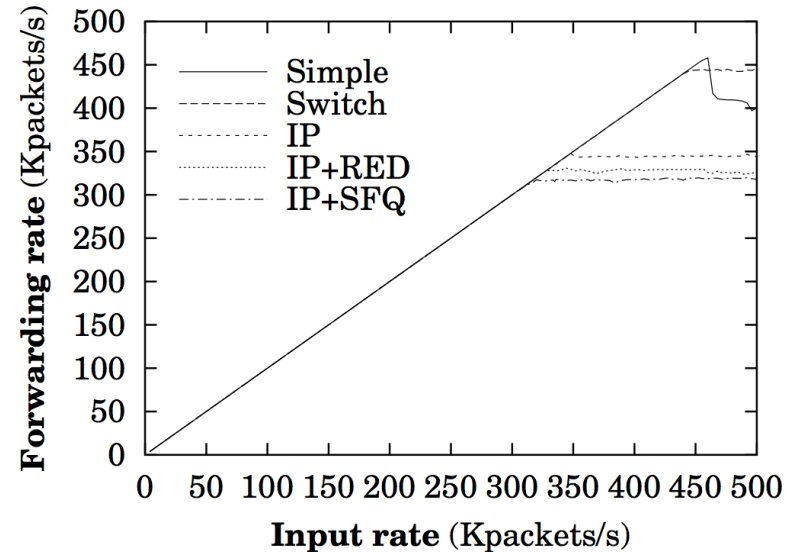
IP router

## IP+RED

Router + random early detection dropping policy

## IP+SFQ

Router + stochastic



Livelock! [Mogul'97]

# Summary

- NFV: everything old is new again
- Cloud networks need
  - Flexible operational model
  - State-of-the-art infrastructure services
  - Programmable network functions
- NFV is changing the networking industry



