



NetBricks: Taking the V out of NFV

Authors: Aurojit Panda, Sangjin Han, Keon Jang, Melvin Walls, Sylvia Ratnasamy, Soctt Shenker

Presented By: 孙晓鹏

2019.5.17



上海交通大学

SHANGHAI JIAO TONG UNIVERSITY

What will cover



- Network functions and their virtualization
- Challenges in virtualization
- Methods invented to overcome these challenges
- NetBricks:
 - Design & Implementation
 - Evaluation

1

Introduction

2

Background & Motivation

3

Design

4

Implementation

5

Evaluation

6

Conclusion



1

Introduction

2

Background & Motivation

3

Design

4

Implementation

5

Evaluation

6


Conclusion



What is NFV?



- Replacing dedicated hardware with software on servers
- Aim: transform network architecture
- No new hardware needed



Why NFV?



- Simplifies adding new functionality: Deploy new software
- Simplifies developing new functionality: Write software vs design hardware
- Reuse management tools form other domains
- Reduce cost

Why NFV unpopular?



- NFV requirements:
 - Performance: latencies & throughput
 - Efficiency: maximize number of NFs on single machine
 - Chaining: each packet needs to be processed by sequence of NFs
- Current tools for building NFs fall short of these requirements:
 - State-of-the-art for NFV is much more primitive
 - VMs/containers incur substantial performance overheads

Contents

1

Introduction

2

Background & Motivation

3

Design

4

Implementation

5

Evaluation

6

Conclusion



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY

Background



- State-of-the-art for NFV is more primitive than that for programming
- Click: does not provide easily customizable low-level optimizations
- DPDK: fast and optimized I/O only
- NFV developers spend much time in code optimization
- More code tweaks may lead to more bugs

Building NFs



- Tools do not support:
 - Rapid development (achieved through high level abstractions)
 - High performance (requiring low-level optimizations)

Building NFs



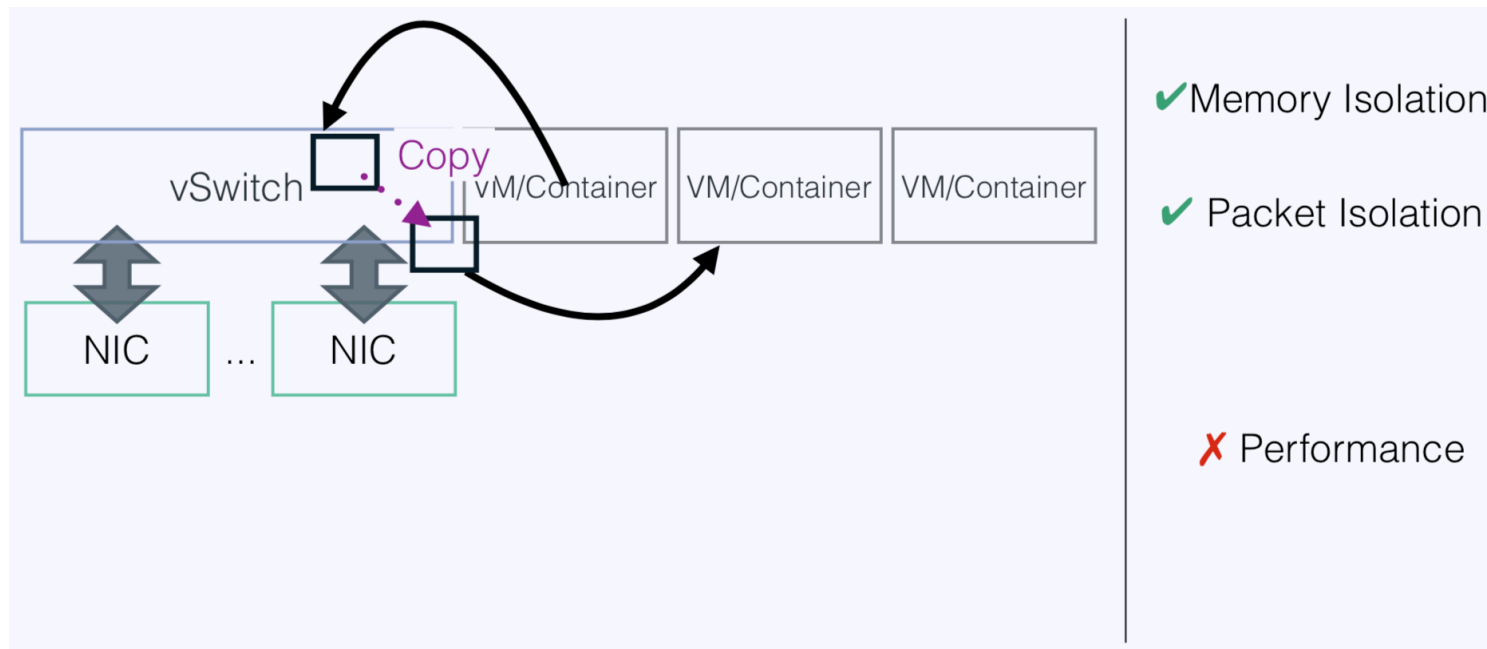
- Click allows NF development by assembling various modules
- Modules support only limited customization
- I/O is optimized, but developers responsible for other optimizations
- Developers often need to implement & optimize new modules

Running NFs



- Isolation between NFs is critical
 - Memory isolation
 - Packet isolation
 - Performance isolation
- Current deployment rely on VMs for isolation
- VMs incur substantial overheads

Current Solution



- During network I/O packets must cross a hardware memory isolation boundary
- This needs a context switch/syscall, which incurs significant overheads

Penalties



Comparison between:

- Single process running a dedicated NIC
- Same functionality on a container
- Same functionality on a VM

Penalties



- For single NF (processing smallest packets - 64B)
 - 3x when using containers
 - 7x when using VMs
- For chained NFs
 - 7x when using containers
 - 11x when using VMs

Contents

1

Introduction

2

Background & Motivation

3

Design

4

Implementation

5

Evaluation

6

Conclusion



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY

Design



- Programming Abstractions
 - Packet Processing
 - Bytestream Processing
 - Control Flow
 - State Abstraction
 - Event Scheduling
- Execution environment
 - Isolation
 - Placement & Scheduling

Abstractions



Packet Processing

Parse/Deparse	Header
Transform	UDF
Filter	UDF

Control Flow

Group By	UDF
Shuffle	UDF
Merge	

Byte Stream

Window	UDF
Packetize	UDF

State

Bounded
Consistency

Packet Processing Abstractions



operation	Input	Process/Output
Parse	Header type and packet structure	Parses the payload using header type and pushes resulting headers onto stack, removes the header bytes from payload
Deparse	-	Pops bottom most header back to payload
Transform	Packet structure, UDF	Modifies header/payload as per UDF
Filter	Packet structure, UDF	Allows packets meeting some criteria (as defined by UDF) to be dropped

Execution environment



- Memory isolation:
 - VM based isolation incurs heavy penalties
 - NetBricks uses software isolation instead
 - Previous work - safe languages with type checks and runtimes can provide memory isolation equivalent to that provided by MMU
 - NetBricks uses Rust (type checking) & LLVM (runtime env)

Execution environment



- Packet Isolation:
 - Usually achieved by copying
 - Zero Copy Soft Isolation (ZCSI)
 - Unique Types - NO simultaneous access to same data from 2 threads
 - Verification at compile time, to avoid runtime overheads

Execution environment



- Placement and Scheduling:
 - NetBricks runs serval NFs
 - NetBricks must decide at compile time what core is to be used to run each NF chain
 - NetBricks must make scheduling decisions about which packet to process next
 - Currently using run-to-completion scheduling
 - Currently using round robin scheduling for deciding event scheduling

Contents

1

Introduction

2

Background & Motivation

3

Design

4

Implementation

5

Evaluation

6

Conclusion



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY

Example NF: Maglev



- Maglev: Load balancer from Google
- NetBricks implementation: 105 lines, 2 hours of time
- Comparable performance to optimized code

```
1 pub fn maglev_nf<T: 'static + NbNode>(  
2     input: T  
3     backends: &[str],  
4     ctx: nb_ctx,  
5     lut_size: usize)  
6     -> Vec<CompositionNode> {  
7     let backend_ct = backends.len();  
8     let lookup_table =  
9         Maglev::new_lut(ctx,  
10             backends,  
11             lut_size);  
12     let mut flow_cache =  
13         BoundedConsistencyMap::<usize, usize>::new();  
14  
15     let groups =  
16         input.shuffle(BuiltInShuffle::flow)  
17             .parse::<MacHeader>()  
18             .group_by(backend_ct, ctx,  
19                 box move |pkt| {  
20                     let hash =  
21                         ipv4_flow_hash(pkt, 0);  
22                     let backend_group =  
23                         flow_cache.entry(hash)  
24                             .or_insert_with(|| {  
25                             lookup_table.lookup(hash) });  
26                     backend_group  
27                 });  
28     groups.iter().map(|g| g.compose()).collect()  
29 }
```

Listing 3: Maglev [9] implemented in NetBricks.

Contents

1

Introduction

2

Background & Motivation

3

Design

4

Implementation

5

Evaluation

6

Conclusion



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY

Setup



- Dual-socket servers equipped with Intel Xeon E5-2660 CPUs
- Each with 10 cores
- Intel XL710QDA2 40Gb NIC
- 2Virtual Switches
- OpenVSwitch with DPDK
- SoftNIC (new virtual switch optimized for NFV use cases)

Overhead for checking array bounds

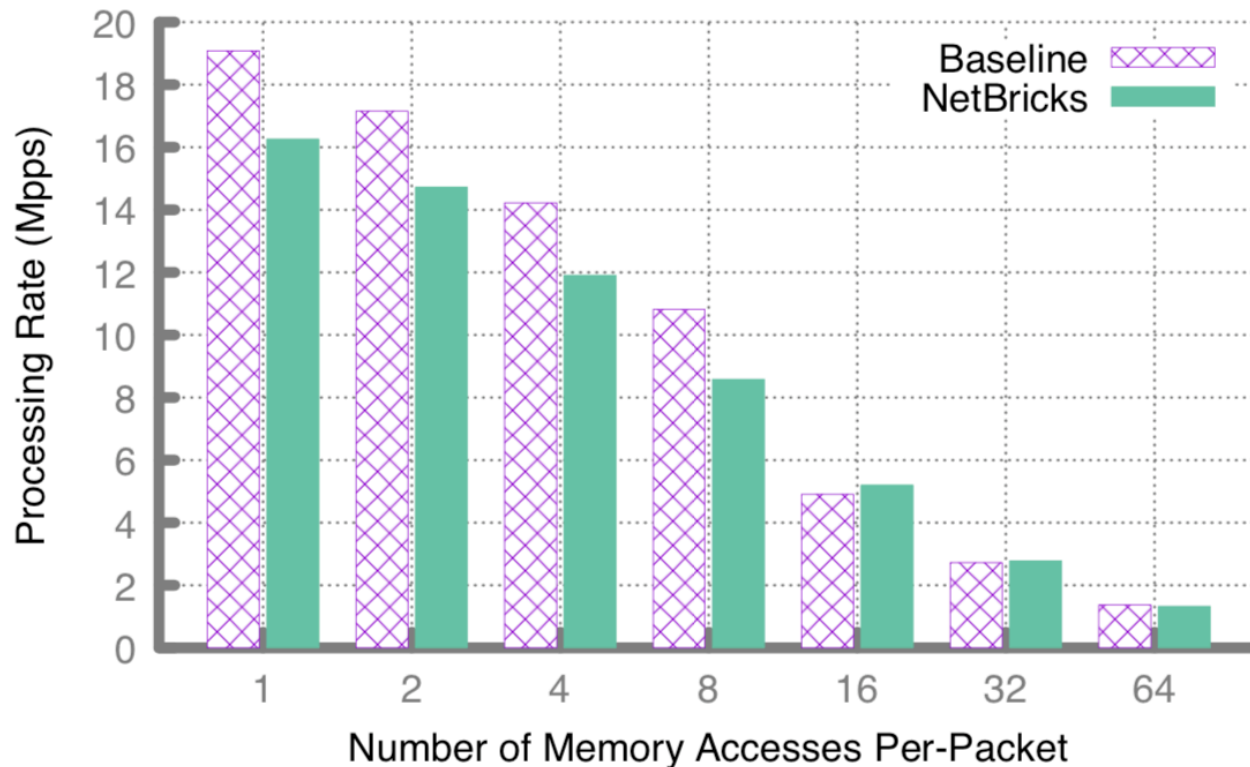


Figure 1: Throughput achieved by a NetBricks NF and an NF written in C using DPDK as the number of memory accesses in a large array grows.

Cost of isolation



- Single NF




Figure 2: Setup for evaluating single NF performance for VMs and containers.




Figure 3: Setup for evaluating single NF performance using NetBricks.

Cost of isolation



- Single NF




Figure 4: Throughput achieved using a single NF running under different isolation environments.

Cost of isolation



■ NF Chains




Figure 5: Setup for evaluating the performance for a chain of NFs, isolated using VMs or Containers.

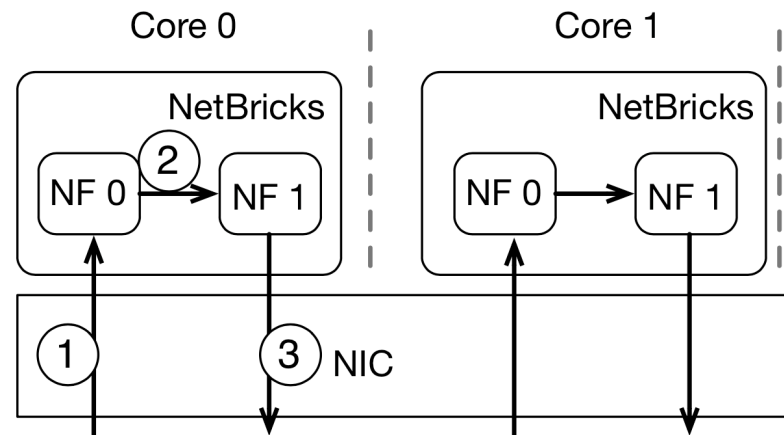


Figure 6: Setup for evaluating the performance for a chaining of NFs, running under NetBricks.

Cost of isolation



■ NF Chains




Figure 7: Throughput with increasing chain length when using 64B packets. In this figure NB-MC represents NetBricks with multiple cores, NB-1C represents NetBricks with 1 core.

Effect of Increasing NF complexity




Figure 9: Throughput for a single NF with increasing number of cycles per-packet using different isolation techniques.

Contents

1

Introduction

2

Background & Motivation

3

Design

4

Implementation

5

Evaluation

6

Conclusion



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY

Conclusion



- High-level programming brings convenience
 - Abstract operators + UDF simplify development
- Software isolation is necessary for high performance NFV
 - Type checking + bound checking + unique types

Related work



- ClickOS (NSDI' 14)
- E2: a framework for NFV applications (SOSP' 15)
- Rollback-Recovery for Middleboxes (SIGCOMM' 15)
- NFP: Enabling Network Function Parallelism in NFV (SIGCOMM' 17)
- Adaptive Interference-Aware VNF Placement for Service-Customized 5G Network Slices (INFOCOM' 19)



Thank you !



Backup

NetBricks Runtime Architecture

