

Bugs: Difference in Beliefs

Rong Chen

Sides from the authors of “Uncertainty to Belief: Inferring the Specification Within” and
“Towards Optimization-Safe Systems Analyzing the Impact of Undefined Behavior”

What is a Bug?

a bug is -

a *contradiction* in beliefs

MUST - implied by the code

a *deviation* from *common behavior*

MAY - inferred from the code

probability of coincidence

Bugs Cost??

Patriot missile defense system

28 **dead** soldiers, 98 wounded

Therac-25 medical device

Several people dead, others wounded

General Electric XA/21

50 million people left without water, electricity.

What Bugs Means to You?



How to find bugs?

What is your belief set?

MUST set

MAY set

What is the implied sets?

Inconsistency means possible bugs!!

Trivial consistency: NULL pointers

*p implies MUST belief:

p is not null

A check (p == NULL) implies two MUST beliefs:

POST: p is null on true path, not null on false path

PRE: p was unknown before check

```
/* 2.4.1: drivers/isdn/svmb1/capidrv.c */  
if (!card)  
    printk(KERN_ERR, "capidrv-%d: ...", card->contrnr...)
```

```
/* drivers/net/wan/sdla_chdlc.c:3948 */  
if (!card) {  
    lock_adapter_irq(&card->wandevice.lock, &smp_flags);  
    card->tty=NULL;
```

Null pointer fun

Use-then-check

```
/* 2.4.7: drivers/char/mxser.c */
struct mxser_struct *info = tty->driver_data;
unsigned flags;
if(!tty || !info->xmit_buf)
    return 0;
```

Contradiction/redundant checks

```
/* 2.4.7/drivers/video/tdfxfb.c */
fb_info.regbase_virt = ioremap_nocache(...);
if(!fb_info.regbase_virt)
    return -ENXIO;
fb_info.bufbase_virt = ioremap_nocache(...);

if(!fb_info.regbase_virt) {
    iounmap(fb_info.regbase_virt);
```

Internal Consistency: finding security holes

Applications are bad:

Rule: “do not dereference user pointer $\langle p \rangle$ ”

One violation = security hole

Big Problem: which are the user pointers???

Sol'n: forall pointers, cross-check two OS beliefs

“ $\ast p$ ” implies safe kernel pointer

“copyin(p)/copyout(p)” implies dangerous user pointer

Error: pointer p has both beliefs.


```

/*
 * Find a routing entry, we only return a FULL match
 */
static struct ipddp_route* ipddp_find_route(struct ipddp_route *rt)
{
    struct ipddp_route *f;

    for(f = ipddp_route_list; f != NULL; f = f->next)
    {
        if(f->ip == rt->ip
            && f->at.s_net == rt->at.s_net
            && f->at.s_node == rt->at.s_node)
            return (f);
    }

    return (NULL);
}

```

- In linux 2.

```

/* drivers/n
case SIOCADD
    return
case SIOCDEL }
    return ipddp_delete(rt);
case SIOFCINDIPDDPRT:
    if(copy_to_user(rt, ipddp_find_route(rt),
        sizeof(struct ipddp_route)))
        return -EFAULT;

```

- “rt” as a tainted pointer, checking warns that rt is passed to a routine that dereferences it

Statistical: Deriving deallocation routines

Use-after free errors are horrible.

Problem: lots of undocumented sub-system free functions

Soln: derive behaviorally: pointer “p” not used after call
“foo(p)” implies **MAY** belief that “foo” is a free function

Conceptually: Assume all functions free all arguments
(in reality: filter functions that have suggestive names)

A bad free error

```
/* drivers/block/cciss.c:cciss_ioctl */
if (ioccommand.Direction == XFER_WRITE) {
    if (copy_to_user(...)) {
        cmd_free(buff, c);
        if (buff != NULL) kfree(buff);
        return(-EFAULT);
    }
}

if (ioccommand.Direction == XFER_READ) {
    if (copy_to_user(...)) {
        cmd_free(buff, c);
        kfree(buff);
    }
}

cmd_free(buff, c);
if (buff != NULL) kfree(buff);
```

“A must be followed by B”

“a(); ... b();” implies MAY belief that a() follows b()

You might believe a-b paired, or might be a coincidence

Checking derived lock functions

Simplest:

```
/* fs/proc/inode.c:41:de_put: */
lock_kernel();
if (!de->count) {
    printk("de_put: entry already free!\n");
    return;
}
unlock_kernel();
```

Evilest:

```
/* 2.4.1: drivers/sound/trident.c:trident_release:
lock_kernel();
card = state->card;
dmabuf = &state->dmabuf;
VALIDATE_STATE(state);
```

```
#define VALIDATE_MAGIC(FOO,MAG)
({
    if (!(FOO) || (FOO)->magic != MAG) {
        printk(invalid magic, FUNCTION);
        return -ENXIO;
    }
})

#define VALIDATE_STATE(a) VALIDATE_MAGIC(a,TRIDENT_STATE_MAGIC)
```

Towards Optimization-Safe Systems

Analyzing the Impact of Undefined Behavior

Xi Wang, Nickolai Zeldovich, M. Frans Kaashoek, Armando Solar-Lezama
MIT CSAIL

Belief: compiler == faithful translator



Not true if your code invokes undefined behavior

- Security implications

Example: compiler discards sanity check

```
char *buf      = ...;
char *buf_end  = ...;
unsigned int off = /* read from untrusted input */;
if (buf + off >= buf_end)
    return;          /* validate off: buf+off too large*/
if (buf + off < buf)
    return;          /* validate off: overflow, buf+off wrapped around */
/* access buf[0..off-1] */
```

- C spec: pointer overflow is undefined behavior
 - gcc: `buf + off` cannot overflow, different from hardware!
 - gcc: `if (buf + off < buf) => if (false)`
- Attack: craft a large off to trigger **buffer overflow**

Undefined behavior allows such optimizations

Undefined behavior: the spec “imposes no requirements”

- Original goal: emit efficient code
- Compilers assume a program **never** invokes undefined behavior
- Example: no bounds checks emitted; assume no buffer overflow

```
*p = 42;           /* store 42 to p */  
    ↓  
mov $42, (%rdi)    /* no bounds checks */
```

Examples of undefined behavior in C

Meaningless checks from real code: pointer p; signed integer x

Pointer overflow: `if (p + 100 < p)`

Signed integer overflow: `if (x + 100 < x)`

Oversized shift: `if (!(1 << x))`

Null pointer dereference: `*p; if (p)`

Absolute value overflow: `if (abs(x) < 0)`

Problem: unstable code confuses programmers

Unstable code: compilers discard code due to undefined behavior



- Security checks discarded
- Weakness amplified
- Unpredictable system behavior



Contributions

- A case study of unstable code in real world
- An algorithm for identifying unstable code
- A static checker STACK
 - 160 previously unknown bugs confirmed and fixed
 - Users: Intel, several open-source projects,...

Example: broken check in Postgres

Implement 64-bit signed division x/y in SQL

```
if (y == -1 && x < 0 && (x / y < 0)) /*  $-2^{63}/-1 < 0$ ? */  
    error();
```

- Some compilers optimize away the check
- x86-64's `idivq` traps on overflow: DoS attack

```
SELECT ((-9223372036854775808)::int8) / (-1);
```

SQL

Example: fix check in Postgres

Our proposal:

```
if (y == -1 && x == INT64_MIN) /* INT64_MIN is  $-2^{63}$ */
```

Developer's fix:

```
if (y == -1 && ((-x < 0) == (x < 0)))
```

- Still unstable code: time bomb for future compilers
 - “it’s an overflow check so it should check for overflow”
 - “we don’t want the constant INT64_MIN; it’s less portable”

“This will create MAJOR SECURITY ISSUES in ALL MANNER OF CODE. I don’t care if your language lawyers tell you gcc is right. . . . **FIX THIS! NOW!**”

a gcc user

bug #30475 - assert(int+100 > int) optimized away

“I am sorry that you wrote broken code to begin with . . . **GCC is not going to change.**”

a gcc developer

bug #30475 - `assert(int+100 > int)` optimized away

Test existing compilers

12 C/C++ compilers

gcc

aCC (HP)

lcc (Intel)

open64 (AMD)

suncc (Oracle)

ti (TI's TMS320C6000)

clang

armcc (ARM)

msvc (Microsoft)

pathcc (PathScale)

xlc(IBM)

windriver (Wind River's Diab)

Examples of unstable code

Meaningless checks from real code: pointer p; signed integer x

Pointer overflow: `if (p + 100 < p)` \Rightarrow `if (false)`

Signed integer overflow: `if (x + 100 < x)` \Rightarrow `if (false)`

Oversized shift: `if (!(1 << x))` \Rightarrow `if (false)`

Null pointer dereference: `*p; if (p)` \Rightarrow `if (false)`

Absolute value overflow: `if (abs(x) < 0)` \Rightarrow `if (false)`

Compilers often discard unstable code

	<code>if(p+100<p)</code>	<code>if(x+100<x)</code>	<code>if(!(1<<x))</code>	<code>*p; if(!p)</code>	<code>if(abs(x)<0)</code>
gcc-4.8.1	O2	O2		O2	O2
clang-3.3	O1	O1	O1		
aCC-6.25					O3
armcc-5.02		O2			
icc-14.0.0		O1		O2	
msvc-14.0.0				O1	
open64-14.0.0	O1	O2			O2
pathcc-1.0.0	O1	O2			O2
suncc-5.12				O3	
ti-7.4.2	O0	O0			
windriver-5.9.2		O0			
xlc-12.1	O3				

Compilers become more aggressive over time

	<code>if(p+100<p)</code>	<code>if(x+100<x)</code>	<code>if(!(1<<x))</code>	<code>*p; if(lp)</code>	<code>if(abs(x)<0)</code>
(1992) gcc-1.42					
(2001) gcc-2.95.3		O1			
(2006) gcc-3.4.6		O1		O2	
(2007) gcc-4.2.1	O0	O2			O2
(2013) gcc-4.8.1	O2	O2		O2	O2
(2009) clang-1.0	O1				
(2010) clang-2.8	O1	O1			
(2013) clang-3.3	O1	O1	O1		

Observation

- Compilers silently remove unstable code
- Different compilers behave in different ways
 - Change/upgrade compiler => broken system
- Need a systematic approach

Our approach: precisely flag unstable code

C/C++ source → LLVM IR → STACK → warnings

```
% ./configure  
% stack-build make          # intercept cc & generate LLVM IR  
% poptck                    # run STACK in parallel
```

STACK provides informative warnings

```
1. res = x / y;  
2. if (y == -1 && x < 0 && res < 0)  
3.     return;
```

The check at line 2 is simplified into false, due to division at line 1

```
model: |                                # possible optimization  
    %cmp3 = icmp slt i64 %res, 0  
    --> false  
stack:                                # location of unstable code  
    - div.c:2  
core:                                  # why optimized away  
    - div.c:1  
    - signed division overflow
```

Design overview of STACK

- What's the difference, compilers vs most programmers?
 - Assumption Δ : programs don't invoke undefined behavior
- What can compilers do only with assumption Δ ?
 - Optimize away unstable code
- STACK: mimic a compiler that selectively enables Δ
 - Phase I: optimize w/o Δ
 - Phase II: optimize w/ Δ
 - Unstable code: difference between the two phases

Example of identifying unstable code

```
1. res = x / y;  
2. if (y == -1 && x < 0 && res < 0)  
3.     return;
```

- Assumption Δ :
 - No division by zero: $y \neq 0$
 - No division overflow: $y \neq -1$ OR $x \neq \text{INT_MIN}$
- STACK can optimize “res < 0” to “false” only with Δ
 - Phase I: is “res < 0” equivalent to “false” in general? No.
 - Phase II: is “res < 0” equivalent to “false” with Δ ? Yes!
- Report “res < 0” as unstable code

Compute assumption Δ

One must not trigger undefined behavior at any code fragment

- **Reach**(e): when to reach and execute code fragment e
- **Undef**(e): when to trigger undefined behavior at e

$$\Delta = \forall e: \text{Reach}(e) \rightarrow \neg \text{Undef}(e)$$

Example: compute assumption Δ

One must not trigger undefined behavior at any code fragment

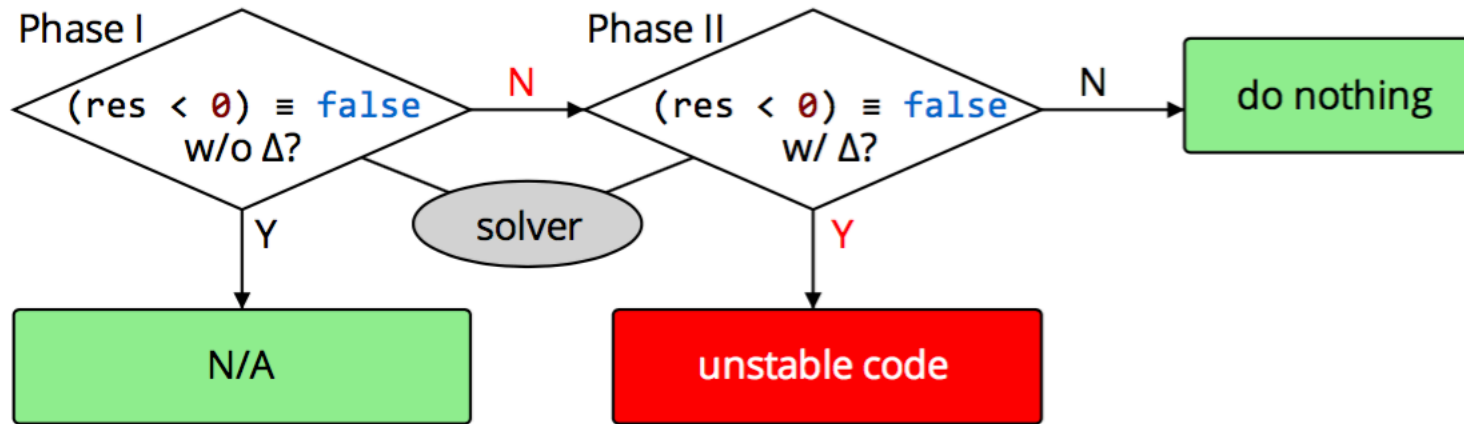
$$\Delta = \forall e: \text{Reach}(e) \rightarrow \neg \text{Undef}(e)$$

```
1. res = x / y;  
2. if (y == -1 && x < 0 && res < 0)  
3.   return;
```

```
 $\Delta = \text{true} \rightarrow \neg((y == 0) \vee (x == -1 \wedge y == \text{INT\_MIN}))$  # line 1  
       $\wedge \text{true} \rightarrow \neg \text{false}$  # line 2  
       $\wedge ((y == -1) \wedge (x < 0) \wedge (x/y < 0)) \rightarrow \neg \text{false}$  # line 3  
  
 $\Delta = \neg((y == 0) \vee (x == -1 \wedge y == \text{INT\_MIN}))$ 
```

Find unstable code by selectively enabling Δ

```
1. res = x / y;  
2. if (y == -1 && x < 0 && res < 0)  
3.   return;
```



Summary of STACK

- Compute assumption Δ : no undefined behavior
- Two-phase framework: w/o and w/ Δ
 - Report unstable code from difference
- Limitations
 - Missing unstable code: Phase II not powerful enough
 - False warnings: Phase I not powerful enough

Implementation of STACK

- LLVM
- Boolector solver
- ~4,000 lines of C++ code
- Per-function for better scalability
 - Could miss bugs

Evaluation

- Is STACK useful for finding unstable code?
- How precise are STACK's warnings?
- How prevalent is unstable code?
- How much time to analyze a large code base?

STACK finds new bugs

- Applied STACK to many popular systems
- Inspected warnings and submitted patches to developers
 - Binutils, Bionic, Dune, e2fsprogs, FFmpeg+Libav, file, FreeType, GMP, GRUB, HiStar, Kerberos, libX11, libarchive, libgcrypt, Linux kernel, Mosh, Mozilla, OpenAFS, OpenSSH, OpenSSL, PHP, plan9port, Postgres, Python, QEMU, Ruby+Rubinius, Sane, uClibc, VLC, Wireshark, Xen, Xpdf
- Developers accepted most of our patches
 - 160 new bugs

STACK warnings are precise

- Kerberos: STACK produced 11 warnings
 - Developers accepted every patch
 - No warnings for fixed code
 - Low false warning rate: 0/11
- Postgres: STACK produced 68 warnings
 - 9 patches accepted: server crash
 - 29 patches in discussion: developers blamed compilers
 - 26 time bombs: can be optimized away by future compilers
 - 4 false warnings: benign redundant code
 - Low false warning rate: 4/68

Unstable code is prevalent

- Applied STACK to all Debian Wheezy packages
 - 8,575 C/C++ packages
 - ~150 days of CPU time to build and analyze
- STACK warns in ~40% of C/C++ packages

STACK scales to large code bases

Intel Core i7-980 3.3 GHz, 6 cores

	build time	analysis time	# files
Kerberos	1 min	2 min	705
Postgres	1 min	11 min	770
Linux kernel	33 min	62 min	14,136

How to avoid unstable code

➤ Programmers

- Fix bugs
- Work around: disable certain optimizations

➤ Compilers & checkers

- Many bug-finding tools fail to model C spec correctly
- Use our ideas to generate better warnings

➤ Language designers: revise the spec

- Eliminate undefined behavior? Perf impact?

Other application

Reflections on trusting trust [Thompson8

➤ Hide backdoors

- Submit a new feature with unstable code
- Could easily slip through code review



Summary

- Compilers optimize away unstable code
 - Subtle bugs
 - Significant security implications
- Compiler writers: use our techniques to generate better warnings
- Language designers: trade-off between performance & security
- Programmers: check your C/C++ code using STACK

<http://css.csail.mit.edu/stack/>